

PyPy 1.4: Status and News

Maciej Fijalkowski

Cape Town PUG

November 27 2010



Outline

- What is PyPy?
- Overview of the PyPy JIT
- cpyext: load CPython extensions in PyPy!

What is PyPy?

- A very compliant Python interpreter
- with a just in time compiler
- ... other features

Speed

- cool charts

1.4 release

- 1.2: released on March 12th, 2010
 - ▶ Main theme: speed
 - ▶ JIT compiler
 - ▶ speed.pypy.org
- 1.3: released on June 26th, 2010
 - ▶ Stability: lot of bugfixes, thanks for the feedback :-)
 - ▶ More speed!
 - ▶ cpyext
- 1.4: release yesterday
 - ▶ even more speed and stability
 - ▶ jitted regexes, 64bit backend

What works on PyPy

- Pure Python modules should Just Work (TM)

- ▶ django trunk
- ▶ twisted, nevow
- ▶ pylons
- ▶ bittorrent
- ▶ ...

- lot of standard modules

- *__builtin__ __pypy__ codecs lsprof minimal curses
random rawffi socket sre weakref bz2 cStringIO crypt errno
exceptions fcntl gc itertools marshal math md5 mmap operator parser
posix pyexpat select sha signal struct symbol sys termios thread time
token unicodedata zipimport zlib*
- *array binascii cPickle cmath collections ctypes datetime functools grp
md5 pwd pyexpat sha sqlite3 syslog*
- *ctypes*

What works on PyPy

- Pure Python modules should Just Work (TM)
 - ▶ django trunk
 - ▶ twisted, nevow
 - ▶ pylons
 - ▶ bittorrent
 - ▶ ...
- lot of standard modules
 - *__builtin__ __pypy__ codecs lsprof minimal curses
random rawffi socket sre weakref bz2 cStringIO crypt errno
exceptions fcntl gc itertools marshal math md5 mmap operator parser
posix pyexpat select sha signal struct symbol sys termios thread time
token unicodedata zipimport zlib*
 - *array binascii cPickle cmath collections ctypes datetime functools grp
md5 pwd pyexpat sha sqlite3 syslog*
 - *ctypes*

What does not work on PyPy

- Pure Python modules should Just Work (TM)
 - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
 - ▶ refcounting: `open('xxx', 'w').write('stuff')`
 - ▶ non-string keys in dict of types (try it!)
 - ▶ exact naming of a list comprehension variable
 - ▶ exact message matching in exception catching code
 - ▶ ...
- Extension modules
 - ▶ try cpyext!

What does not work on PyPy

- Pure Python modules should Just Work (TM)
 - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
 - ▶ refcounting: `open('xxx', 'w').write('stuff')`
 - ▶ non-string keys in dict of types (try it!)
 - ▶ exact naming of a list comprehension variable
 - ▶ exact message matching in exception catching code
 - ▶ ...
- Extension modules
 - ▶ try cpyext!

What does not work on PyPy

- Pure Python modules should Just Work (TM)
 - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
 - ▶ refcounting: `open('xxx', 'w').write('stuff')`
 - ▶ non-string keys in dict of types (try it!)
 - ▶ exact naming of a list comprehension variable
 - ▶ exact message matching in exception catching code
 - ▶ ...
- Extension modules
 - ▶ try cpyext!

What does not work on PyPy

- Pure Python modules should Just Work (TM)
 - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
 - ▶ refcounting: `open('xxx', 'w').write('stuff')`
 - ▶ non-string keys in dict of types (try it!)
 - ▶ exact naming of a list comprehension variable
 - ▶ exact message matching in exception catching code
 - ▶ ...
- Extension modules
 - ▶ try cpyext!

Mandelbrot demo XXX maybe

Part 2: Just-in-Time compilation

Snakes never crawled so fast

Overview of implementations

- CPython
- Stackless
- Psyco
- Jython
- IronPython
- PyPy

Features

- it just works
- it may give good speed-ups (better than Psyco)
- it may have a few bugs left (Psyco too)
- it is not a hack (unlike Psyco)
- PyPy also has excellent memory usage
 - ▶ half that of CPython for a program using several hundreds MBs

Features

- it just works
- it may give good speed-ups (better than Psyco)
- it may have a few bugs left (Psyco too)
- it is not a hack (unlike Psyco)
- PyPy also has excellent memory usage
 - ▶ half that of CPython for a program using several hundreds MBs

Features

- it just works
- it may give good speed-ups (better than Psyco)
- it may have a few bugs left (Psyco too)
- it is not a hack (unlike Psyco)
- PyPy also has excellent memory usage
 - ▶ half that of CPython for a program using several hundreds MBs

What is a JIT

- CPython compiles the program source into bytecodes
- without a JIT, the bytecodes are then interpreted
- with a JIT, the bytecodes are further translated to machine code (assembler)

What is a JIT (2)

The translation can be:

- syntactic: translate the whole functions into machine code
 - ▶ “the obvious way”
 - ▶ e.g. Pyrex/Cython
 - ▶ not good performance, or needs tricks
- semantic: translate bits of the function just-in-time
 - ▶ only used parts
 - ▶ exploit runtime information (e.g. types)
 - ▶ Psyco, PyPy

What is a JIT (2)

The translation can be:

- syntactic: translate the whole functions into machine code
 - ▶ “the obvious way”
 - ▶ e.g. Pyrex/Cython
 - ▶ not good performance, or needs tricks
- semantic: translate bits of the function just-in-time
 - ▶ only used parts
 - ▶ exploit runtime information (e.g. types)
 - ▶ Psyco, PyPy

What is a mixed JIT

- start by interpreting normally
- find loops as they are executed
- turn them into machine code
- 80% of the time is spent in 20% of the code

Speed of the PyPy JIT

- most programs work quite well
- dynamizm that can't be reduced is not good
- meta-programming to the rescue
- reading frames or setting trace hooks has a bad effect
- generators are slower than they should be
- pure python is generally better than lower-level loops (like itertools)

Optimizations part 1 - no frames

- frame access is delayed
- if we want frames (`sys._getframe`, `sys.exc_info`, ...), it works by reading out of processor stack

Optimizations part 2 - malloc removal

- in python each object (including `int`) has to be allocated
- if we can prove that the object does not escape, we can ignore allocation
- in case we exit the JIT, we can create it on the fly
- much more powerful than proving compilation time

Optimizations part 3 - inlining

- very powerful optimization
- python calling is insane
- if we remove everything, we can even not allocate the frame
- huge wins

Optimizations part 4 - map dicts

- avoiding dict lookups
- work as good as `__slots__` memory-wise
- much faster when jitted

Part 3

cpyext

cpyext

- CPython extension modules in PyPy
- `pypy-c setup.py build`
- included in PyPy 1.3
- still beta
- 50% of the CPython API is supported
 - ▶ enough for 90% of extension modules

features

- C API written in Python!
- Testable on top of an interpreted py.py
- Written on top of the object space
- Source compatibility
 - ▶ PyString_AS_STRING is actually a function call (instead of a macro)

```
@cpython_api([PyObject], Py_ssize_t, error=-1)
def PyDict_Size(space, w_obj):
    return space.int_w(space.len(w_obj))
```

implementation

- It was not supposed to work!
 - ▶ different garbage collector
 - ▶ no “borrowed reference”
 - ▶ all the PyObject slots
- not faster than python code!
- PyObject contains ob_type and ob_refcnt
 - ▶ The “abstract object interface” is used.
- Some objects contain more:
 - ▶ PyString_AsString() must keep the buffer alive at a fixed location
 - ▶ PyObject exposes all its fields

The Reference Counting Issue

- pypy uses a moving garbage collector, starts with static roots to find objects.
- CPython objects don't move, and PyObject* can point to deallocated memory.
- cpyext builds PyObject as proxies to the “real” interpreter objects
- one dictionary lookup each time the boundary is crossed
- More tricks needed for borrowing references
 - ▶ The object lifetime is tied to its container.
 - ▶ “out of nothing” borrowed references are kept until the end of the current pypy->C call.

supported modules

- Known to work (after small patches):
 - ▶ wxPython
 - ▶ `_sre`
 - ▶ PyCrypto
 - ▶ PIL
 - ▶ cx_Oracle
 - ▶ MySQLdb
 - ▶ sqlite

Why your module will crash


Likely:

```
static PyObject *myException;

void init_foo() {
    myException = PyErr_NewException(...);
    Py_AddModule(m, myException); // steals a reference
}

{
    PyErr_SetString(myException, "message"); // crash
}
```

wxPython on PyPy (1)



wxPython on PyPy (2)

The screenshot shows the wxPython: (A Demonstration) application window. The left sidebar contains a tree view of wxPython demos, with 'I18N' selected. The main area displays an error message: 'An error has occurred while trying to run the demo'. The exception is of type '<type 'exceptions.AttributeError'>' with details: 'module' object has no attribute 'getrefcount'. Below the error message is a traceback table.

File...	Line	Function	Code
Main.py	1901	RunModule	self.demoPage = module.runTest(self, self.nb, self)
I18N.py	214	runTest	win = LanguageSelectPanel(nb, log)
I18N.py	165	__init__	self.OnLangSelectAndTranslate()
I18N.py	202	OnLangSelectAndTranslate	self.updateLanguage(lang)
I18N.py	175	updateLanguage	assert sys.getrefcount(self.locale) <= 2

Below the traceback table, a note states: 'Entries from the demo module are shown in blue. Double-click on them to go to the offending line.' At the bottom, the 'Demo Log Messages' section shows a list of log entries, including 'OnActivate: True', 'OnActivate: False', and 'Loading demo RendererNative.py...'. The 'Filter Demos' section at the bottom left has a search box.

Contact / Q&A

- Maciej Fijalkowski: fijall at gmail
- The #pypy IRC channel on freenode.net!
- Links:
 - ▶ PyPy: <http://pypy.org/>
 - ▶ PyPy speed center: <http://speed.pypy.org/>
 - ▶ Blog: <http://morepypy.blogspot.com>