

Loop-Aware Optimizations in PyPy's Tracing JIT

Håkan Ardö¹ Carl Friedrich Bolz² Maciej Fijałkowski

¹Centre for Mathematical Sciences, Lund University

²Heinrich-Heine-Universität Düsseldorf, STUPS Group, Germany

2012 DLS, 22nd of October, 2012

- Context: RPython
- a language for writing interpreters for dynamic languages
- a generic tracing JIT, applicable to many languages
- used to implement PyPy, an efficient Python interpreter

Tracing JITs Compile by Observing an Interpreter

- VM contains both an interpreter and the tracing JIT compiler
- JIT works by observing and logging what the interpreter does
- for interesting, commonly executed code paths
- produces a linear list of operations (trace)

Why do tracing JITs work?

- They are good at selecting interesting and common code paths
- both through the user program and through the runtime
- the latter is particularly important for dynamic languages with a big runtime like Python

Why do tracing JITs work?

- They are good at selecting interesting and common code paths
- both through the user program and through the runtime
- the latter is particularly important for dynamic languages with a big runtime like Python
- traces are trivial to optimize

Optimizing traces

- A trace is an extended basic block
- (it has one entry and several exits)
- traces are easy to optimize due to lack of control flow merges

Optimizing traces

- A trace is an extended basic block
- (it has one entry and several exits)
- traces are easy to optimize due to lack of control flow merges
- most optimizations are one forward pass
- optimizers are often like symbolic executors
- can do optimizations that are expensive or even untractable with full control flow

Example



```
 $w_0 = a_0 + 0$   
 $x_0 = w_0 + 1$   
 $y_0 = a_0 + 1$   
 $z_0 = x_0 + y_0$   
...  
 $\text{jump}(L_0, z_0, b_0)$ 
```



Example



~~$w_0 = a_0 + 0$~~

$x_0 = w_0 + 1$

$y_0 = a_0 + 1$

$z_0 = x_0 + y_0$

...

$\text{jump}(L_0, z_0, b_0)$

rename w_0 to a_0

Example



~~$w_0 = a_0 + 0$~~

$x_0 = a_0 + 1$

$y_0 = a_0 + 1$

$z_0 = x_0 + y_0$

...

$\text{jump}(L_0, z_0, b_0)$

rename w_0 to a_0

$a_0 + 1$ in x_0

Example

~~$w_0 = a_0 + 0$~~

$x_0 = a_0 + 1$

~~$y_0 = a_0 + 1$~~

$z_0 = x_0 + y_0$

...

$\text{jump}(L_0, z_0, b_0)$

rename w_0 to a_0

$a_0 + 1$ in x_0

rename y_0 to x_0

Example

~~$w_0 = a_0 + 0$~~

$x_0 = a_0 + 1$

~~$y_0 = a_0 + 1$~~

$z_0 = x_0 + x_0$

...

jump(L_0 , z_0 , b_0)

rename w_0 to a_0

$a_0 + 1$ in x_0

rename y_0 to x_0

$x_0 + x_0$ in z_0

Problems with this approach

- most traces actually are loops
- naive forward passes ignore this bit of control flow optimization available
- how to fix that without sacrificing simplicity of optimizations?

Idea for solution

- idea first proposed and implemented in LuaJIT by Mike Pall
- this talk presents the implementation of the same approach in RPython's tracing JIT

- idea first proposed and implemented in LuaJIT by Mike Pall
- this talk presents the implementation of the same approach in RPython's tracing JIT

Approach

- do a pre-processing step on the traces
- apply the unchanged forward-pass optimizations
- do some post-processing
- pre-processing is done in such a way that the normal optimizations become loop-aware

- idea first proposed and implemented in LuaJIT by Mike Pall
- this talk presents the implementation of the same approach in RPython's tracing JIT

Approach


- do a pre-processing step on the traces
- apply the unchanged forward-pass optimizations
- do some post-processing
- pre-processing is done in such a way that the normal optimizations become loop-aware
- intuition: give the optimizations a second iteration of context to work with

Pre-processing the loops

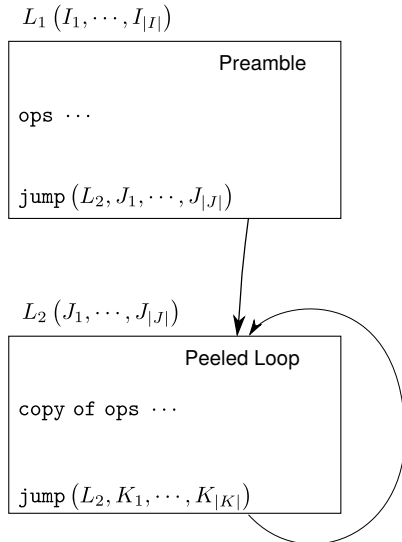
- pre-processing does loop unrolling
- peels off one iteration of the loop, duplicating the trace
- the optimizations optimize both iterations together
- this yields loop-invariant code motion and related optimizations

Loop Peeling

Original Loop:



After Loop Peeling:



Loop Peeling

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_0, i_0$ )
```

Loop Peeling

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_0, i_0$ )
```

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

```
 $L_1(i_0):$   
 $i_2 = i_0 + 1$   
print( $i_2$ )  
jump( $L_1, i_0$ )
```

Apply Optimizations

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

```
 $L_1(i_0):$   
 $i_2 = i_0 + 1$   
print( $i_2$ )  
jump( $L_1, i_0$ )
```

Apply Optimizations

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

```
 $L_1(i_0):$   
 $i_2 = i_0 + 1$   
print( $i_2$ )  
jump( $L_1, i_0$ )
```

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

```
 $L_1(i_0):$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

Add extra arguments

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

```
 $L_1(i_0):$   
print( $i_1$ )  
jump( $L_1, i_0$ )
```

Add extra arguments

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1$ ,  $i_0$ )
```

```
 $L_1(i_0):$   
print( $i_1$ )  
jump( $L_1$ ,  $i_0$ )
```

```
 $L_0(i_0):$   
 $i_1 = i_0 + 1$   
print( $i_1$ )  
jump( $L_1$ ,  $i_0$ ,  $i_1$ )
```

```
 $L_1(i_0, i_1):$   
print( $i_1$ )  
jump( $L_1$ ,  $i_0$ ,  $i_1$ )
```


Optimizations helped by loop peeling

- redundant guard removal
- common subexpression elimination
- heap optimizations
- allocation removal

Larger example

```
while True:  
    y = y + 1
```

Larger example

```
while True:  
    y = y + 1
```

```
 $L_0(p_0, p_1)$ :  
guard_class( $p_1$ , BoxedInteger)  
 $i_2 = \text{get}(p_1, \text{intval})$   
guard_class( $p_0$ , BoxedInteger)  
 $i_3 = \text{get}(p_0, \text{intval})$   
 $i_4 = i_2 + i_3$   
 $p_5 = \text{new}(\text{BoxedInteger})$   
 $\text{set}(p_5, \text{intval}, i_4)$   
 $\text{jump}(L_0, p_0, p_5)$ 
```

Peeled trace

```
 $L_0(p_0, p_1)$ :  
guard_class( $p_1$ , BoxedInteger)  
 $i_2 = \text{get}(p_1, \text{intval})$   
guard_class( $p_0$ , BoxedInteger)  
 $i_3 = \text{get}(p_0, \text{intval})$   
 $i_4 = i_2 + i_3$   
 $p_5 = \text{new}(\text{BoxedInteger})$   
set( $p_5$ , intval,  $i_4$ )  
jump( $L_1$ ,  $p_0$ ,  $p_5$ )
```

```
 $L_1(p_0, p_5)$ :  
guard_class( $p_5$ , BoxedInteger)  
 $i_6 = \text{get}(p_5, \text{intval})$   
guard_class( $p_0$ , BoxedInteger)  
 $i_7 = \text{get}(p_0, \text{intval})$   
 $i_8 = i_6 + i_7$   
 $p_9 = \text{new}(\text{BoxedInteger})$   
set( $p_9$ , intval,  $i_8$ )  
jump( $L_1$ ,  $p_0$ ,  $p_9$ )
```

Final trace


```
 $L_0(p_0, p_1)$ :  
guard_class( $p_1$ , BoxedInteger)  
 $i_2 = \text{get}(p_1, \text{intval})$   
guard_class( $p_0$ , BoxedInteger)  
 $i_3 = \text{get}(p_0, \text{intval})$   
 $i_4 = i_2 + i_3$   
jump( $L_1, p_0, i_4$ )
```

```
 $L_1(p_0, i_3, i_4)$ :  
 $i_8 = i_4 + i_3$   
jump( $L_1, p_0, i_3, i_8$ )
```

- a number of numeric kernels
- some for image processing
- some from SciMark
- comparison against GCC and LuaJIT

- a number of numeric kernels
- some for image processing
- some from SciMark
- comparison against GCC and LuaJIT
- geometric mean of speedups of loop peeling is 70%

Benchmark Results



- a simple preprocessing step on traces enables loop-aware optimizations for tracing JITs
- only minimal changes to the existing optimizations necessary

- Video analytics research example
 - Experimenten driven - prototyping
 - Custom loops over the pixels
 - Good enough performace
-
- Image class with task specific features
 - Zero-padded
 - Clips updates outside border
 - @autoreload decorator reloading functions on code change
 - ReloadHack class reloads and reinstanciates on code change

Image class

```
class Image(object):
    def __getitem__(self, (x, y)):
        if 0 <= x < self.width and 0 <= y < self.height:
            return self.data[y * self.width + x]
        return 0

    def __setitem__(self, (x, y), value):
        if 0 <= x < self.width and 0 <= y < self.height:
            self.data[y * self.width + x] = value

    __add__ = binop(float.__add__)
    __sub__ = binop(float.__sub__)
    __mul__ = binop(float.__mul__)
    __div__ = binop(float.__div__)
    __pow__ = binop(float.__pow__)

    ...
```

Image class

```
def binop(op):
    def f(a, b):
        if not isinstance(a, Image):
            a = ConstantImage(b.width, b.height, a)
        if not isinstance(b, Image):
            b = ConstantImage(a.width, a.height, b)

        out = a.new(typecode='d')
        for x, y in a.indexes():
            out[x, y] = op(float(a[x, y]), float(b[x, y]))

        return out
    return f
```