

# PyPy 1.2 1.3: Status and News

Amaury Forgeot d'Arc  
Antonio Cuni  
Armin Rigo

EuroPython 2010

July 19 2010



# Outline

- PyPy 1.3: what's new and status update
- Overview of the JIT
- cpyext: load CPython extensions in PyPy!

# Part 0: What is PyPy? :-)

- Python interpreter written in Python
  - Framework for developing dynamic languages
  - etc. etc.
- 
- From the user point of view
  - An alternative to CPython
  - with more features!

# Part 0: What is PyPy? :-)

- Python interpreter written in Python
  - Framework for developing dynamic languages
  - etc. etc.
- 
- From the user point of view
  - An alternative to CPython
  - with more features!

# Part 1

- What's new and status update

# What's new in PyPy 1.2, 1.3

- 1.2: released on March 12th, 2010
  - ▶ Main theme: speed
  - ▶ JIT compiler
  - ▶ [speed.pypy.org](http://speed.pypy.org)
- 1.3: released on June 26th, 2010
  - ▶ Stability: lot of bugfixes, thanks for the feedback :-)
  - ▶ More speed!
  - ▶ cpyext
- Binaries for Linux, Windows, Mac
- Ubuntu packages

# Speed: PyPy vs CPython

| Benchmark          | Time (s) | std dev | Current change | Trend  | Times | (log2 scale) |    |       |    |    |    |     |  |
|--------------------|----------|---------|----------------|--------|-------|--------------|----|-------|----|----|----|-----|--|
|                    |          |         |                |        |       | 4x           | 2x | equal | 2x | 4x | 8x | 16x |  |
| richards           | 0.0246   | 0.0018  | 0.32%          | -0.23% | 12.04 |              |    |       |    |    |    |     |  |
| spectral-norm      | 0.0417   | 0.0060  | -4.43%         | -3.90% | 11.57 |              |    |       |    |    |    |     |  |
| chaos              | 0.0426   | 0.0327  | 4.59%          | 1.31%  | 9.77  |              |    |       |    |    |    |     |  |
| nbody_modified     | 0.0810   | 0.0086  | 2.48%          | 0.44%  | 6.63  |              |    |       |    |    |    |     |  |
| float              | 0.0895   | 0.0285  | -1.83%         | -2.37% | 5.64  |              |    |       |    |    |    |     |  |
| fannkuch           | 0.3754   | 0.0122  | 1.64%          | 0.60%  | 4.86  |              |    |       |    |    |    |     |  |
| twisted_iteration  | 0.0329   | 0.0003  | 0.10%          | -0.64% | 4.02  |              |    |       |    |    |    |     |  |
| telco              | 0.3382   | 0.0155  | -0.41%         | -1.13% | 2.95  |              |    |       |    |    |    |     |  |
| django             | 0.2834   | 0.0071  | -0.48%         | -1.05% | 2.85  |              |    |       |    |    |    |     |  |
| spitfire_cstringio | 4.1262   | 0.1360  | -7.40%         | -1.18% | 2.06  |              |    |       |    |    |    |     |  |
| twisted_pb         | 0.0313   | 0.0003  | -0.67%         | -2.43% | 1.83  |              |    |       |    |    |    |     |  |
| twisted_names      | 0.0063   | 0.0000  | -0.68%         | -1.19% | 1.41  |              |    |       |    |    |    |     |  |
| html5lib           | 10.1353  | 1.5026  | 2.37%          | 3.04%  | 1.15  |              |    |       |    |    |    |     |  |
| twisted_web        | 0.1008   | 0.0079  | -0.39%         | -1.04% | 1.12  |              |    |       |    |    |    |     |  |
| rietveld           | 0.4946   | 0.2217  | -0.87%         | 0.21%  | 0.96  |              |    |       |    |    |    |     |  |
| spitfire           | 7.1760   | 0.1540  | -0.82%         | 2.35%  | 0.96  |              |    |       |    |    |    |     |  |
| ai                 | 0.4504   | 0.0078  | -3.11%         | 1.08%  | 0.94  |              |    |       |    |    |    |     |  |
| meteor-contest     | 0.3509   | 0.0135  | -2.62%         | -2.47% | 0.88  |              |    |       |    |    |    |     |  |
| twisted_tcp        | 1.6573   | 0.0176  | -2.88%         | -1.52% | 0.57  |              |    |       |    |    |    |     |  |
| slowspitfire       | 1.2217   | 0.1532  | -7.55%         | -0.77% | 0.56  |              |    |       |    |    |    |     |  |
| spambayes          | 0.8845   | 0.0646  | -3.65%         | -0.24% | 0.31  |              |    |       |    |    |    |     |  |

# Speed: PyPy vs Psyco

| Benchmark          | Time (s) | std dev | Current change | Trend  | Times | (log2 scale) |    |       |    |    |    |     |
|--------------------|----------|---------|----------------|--------|-------|--------------|----|-------|----|----|----|-----|
|                    |          |         |                |        |       | 4x           | 2x | equal | 2x | 4x | 8x | 16x |
| twisted_iteration  | 0.0329   | 0.0003  | 0.10%          | -0.64% | 3.71  |              |    |       |    |    |    |     |
| chaos              | 0.0426   | 0.0327  | 4.59%          | 1.31%  | 3.67  |              |    |       |    |    |    |     |
| django             | 0.2834   | 0.0071  | -0.48%         | -1.05% | 2.79  |              |    |       |    |    |    |     |
| nbody_modified     | 0.0810   | 0.0086  | 2.48%          | 0.44%  | 2.28  |              |    |       |    |    |    |     |
| float              | 0.0895   | 0.0285  | -1.83%         | -2.37% | 2.23  |              |    |       |    |    |    |     |
| spectral-norm      | 0.0417   | 0.0060  | -4.43%         | -3.90% | 2.17  |              |    |       |    |    |    |     |
| richards           | 0.0246   | 0.0018  | 0.32%          | -0.23% | 2.15  |              |    |       |    |    |    |     |
| fannkuch           | 0.3754   | 0.0122  | 1.64%          | 0.60%  | 1.84  |              |    |       |    |    |    |     |
| telco              | 0.3382   | 0.0155  | -0.41%         | -1.13% | 1.70  |              |    |       |    |    |    |     |
| html5lib           | 10.1353  | 1.5026  | 2.37%          | 3.04%  | 1.59  |              |    |       |    |    |    |     |
| twisted_names      | 0.0063   | 0.0000  | -0.68%         | -1.19% | 1.55  |              |    |       |    |    |    |     |
| twisted_web        | 0.1008   | 0.0079  | -0.39%         | -1.04% | 1.33  |              |    |       |    |    |    |     |
| ai                 | 0.4504   | 0.0078  | -3.11%         | 1.08%  | 1.31  |              |    |       |    |    |    |     |
| rietveld           | 0.4946   | 0.2217  | -0.87%         | 0.21%  | 1.29  |              |    |       |    |    |    |     |
| spitfire_cstringio | 4.1262   | 0.1360  | -7.40%         | -1.18% | 1.27  |              |    |       |    |    |    |     |
| twisted_pb         | 0.0313   | 0.0003  | -0.67%         | -2.43% | 1.00  |              |    |       |    |    |    |     |
| spitfire           | 7.1760   | 0.1540  | -0.82%         | 2.35%  | 0.59  |              |    |       |    |    |    |     |
| twisted_tcp        | 1.6573   | 0.0176  | -2.88%         | -1.52% | 0.59  |              |    |       |    |    |    |     |
| meteor-contest     | 0.3509   | 0.0135  | -2.62%         | -2.47% | 0.49  |              |    |       |    |    |    |     |
| slowspitfire       | 1.2217   | 0.1532  | -7.55%         | -0.77% | 0.38  |              |    |       |    |    |    |     |
| spambayes          | 0.8845   | 0.0646  | -3.65%         | -0.24% | 0.34  |              |    |       |    |    |    |     |



# What works on PyPy

- Pure Python modules should Just Work (TM)

- ▶ django trunk
- ▶ twisted, nevow
- ▶ pylons
- ▶ bittorrent
- ▶ ...

- lot of standard modules

- *\_\_builtin\_\_ \_\_pypy\_\_ codecs lsprof minimal curses  
\_\_random\_\_ rawffi socket sre weakref bz2 cStringIO crypt errno  
exceptions fcntl gc itertools marshal math md5 mmap operator parser  
posix pyexpat select sha signal struct symbol sys termios thread time  
token unicodedata zipimport zlib*
- *array binascii cPickle cmath collections ctypes datetime functools grp  
md5 pwd pyexpat sha sqlite3 syslog*
- *ctypes*

# What works on PyPy

- Pure Python modules should Just Work (TM)
  - ▶ django trunk
  - ▶ twisted, nevow
  - ▶ pylons
  - ▶ bittorrent
  - ▶ ...
- lot of standard modules
  - *\_\_builtin\_\_ \_\_pypy\_\_ codecs lsprof minimal curses  
random rawffi socket sre weakref bz2 cStringIO crypt errno  
exceptions fcntl gc itertools marshal math md5 mmap operator parser  
posix pyexpat select sha signal struct symbol sys termios thread time  
token unicodedata zipimport zlib*
  - *array binascii cPickle cmath collections ctypes datetime functools grp  
md5 pwd pyexpat sha sqlite3 syslog*
  - *ctypes*

# What does not work on PyPy

- Pure Python modules should Just Work (TM)
  - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
  - ▶ refcounting: `open('xxx', 'w').write('stuff')`
  - ▶ non-string keys in dict of types (try it!)
  - ▶ exact naming of a list comprehension variable
  - ▶ exact message matching in exception catching code
  - ▶ ...
- Extension modules
  - ▶ try cpyext!

# What does not work on PyPy

- Pure Python modules should Just Work (TM)
  - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
  - ▶ refcounting: `open('xxx', 'w').write('stuff')`
  - ▶ non-string keys in dict of types (try it!)
  - ▶ exact naming of a list comprehension variable
  - ▶ exact message matching in exception catching code
  - ▶ ...
- Extension modules
  - ▶ try cpyext!

# What does not work on PyPy

- Pure Python modules should Just Work (TM)
  - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
  - ▶ refcounting: `open('xxx', 'w').write('stuff')`
  - ▶ non-string keys in dict of types (try it!)
  - ▶ exact naming of a list comprehension variable
  - ▶ exact message matching in exception catching code
  - ▶ ...
- Extension modules
  - ▶ try cpyext!

# What does not work on PyPy

- Pure Python modules should Just Work (TM)
  - ▶ ... unless they don't :-)
- Programs that rely on CPython-specific behavior
  - ▶ refcounting: `open('xxx', 'w').write('stuff')`
  - ▶ non-string keys in dict of types (try it!)
  - ▶ exact naming of a list comprehension variable
  - ▶ exact message matching in exception catching code
  - ▶ ...
- Extension modules
  - ▶ try cpyext!

# Speed: Demo

- Django application
- Mandelbrot fractal
  - ▶ fished randomly on the net :-)
- Run both on CPython and PyPy
  - ▶ django trunk!

# Mandelbrot demo

- Works purely on PyPy
- Not always the case
  - ▶ missing extension modules (cpyext mitigates the problem)
  - ▶ libraries that rely on CPython details
  - ▶ ...
- clear performance-critical part



# CPython and PyPy side by side

- CPython: runs the main application
- PyPy: subprocess, runs only the hotspots
- How do they communicate?
- execnet
  - ▶ **The Ring of Python**, Holger Krekel, 9:45
  - ▶ ouch, too late :-)

# Rendering (1)

## Mandelbrot

```
def render(request):  
    w = int(request.GET.get('w', 320))  
    h = int(request.GET.get('h', 240))  
  
    from py_mandel import mandelbrot  
    img = mandelbrot(w, h)  
  
    return HttpResponse(img, content_type="image/bmp")
```

# Rendering (2)

## Mandelbrot on PyPy

```
def pypy_render(request):  
    w = int(request.GET.get('w', 320))  
    h = int(request.GET.get('h', 240))  
  
    channel = pypy.remote_exec("""  
        from py_mandel import mandelbrot  
        w, h = channel.receive()  
        img = mandelbrot(w, h)  
        channel.send(img)  
    """)  
    channel.send((w, h))  
    img = channel.receive()  
  
    return HttpResponse(img, content_type="image/bmp")
```


# execnet setup

## At startup

```
import execnet
mygroup = execnet.Group()
pypy = mygroup.makegateway("popen//python=pypy-c")
pypy.remote_exec("""
    import sys
    import os
    os.chdir("mandelbrot")
    sys.path.insert(0, '')
""")
```

# Demo

# Benchmarks



## Part 2: Just-in-Time compilation

*Snakes never crawled so fast*

# Overview of implementations

- CPython
- Stackless
- Psyco
- Jython
- IronPython
- PyPy (without and with JIT)
- Unladen Swallow



# Demo

# Features

- it just works
- it may give good speed-ups (better than Psyco)
- it may have a few bugs left (Psyco too)
- it is not a hack (unlike Psyco)
- PyPy also has excellent memory usage
  - ▶ half that of CPython for a program using several hundreds MBs

# Features

- it just works
- it may give good speed-ups (better than Psyco)
- it may have a few bugs left (Psyco too)
- it is not a hack (unlike Psyco)
- PyPy also has excellent memory usage
  - ▶ half that of CPython for a program using several hundreds MBs

# Features

- it just works
- it may give good speed-ups (better than Psyco)
- it may have a few bugs left (Psyco too)
- it is not a hack (unlike Psyco)
- PyPy also has excellent memory usage
  - ▶ half that of CPython for a program using several hundreds MBs

# What is a JIT

- CPython compiles the program source into bytecodes
- without a JIT, the bytecodes are then interpreted
- with a JIT, the bytecodes are further translated to machine code (assembler)

# What is a JIT (2)

The translation can be:

- syntactic: translate the whole functions into machine code
  - ▶ “the obvious way”
  - ▶ e.g. Pyrex/Cython, Unladen Swallow
  - ▶ not good performance, or needs tricks
- semantic: translate bits of the function just-in-time
  - ▶ only used parts
  - ▶ exploit runtime information (e.g. types)
  - ▶ Psyco, PyPy

# What is a JIT (2)

The translation can be:

- syntactic: translate the whole functions into machine code
  - ▶ “the obvious way”
  - ▶ e.g. Pyrex/Cython, Unladen Swallow
  - ▶ not good performance, or needs tricks
- semantic: translate bits of the function just-in-time
  - ▶ only used parts
  - ▶ exploit runtime information (e.g. types)
  - ▶ Psyco, PyPy

# What is a tracing JIT

- start by interpreting normally
- find loops as they are executed
- turn them into machine code
- 80% of the time is spent in 20% of the code



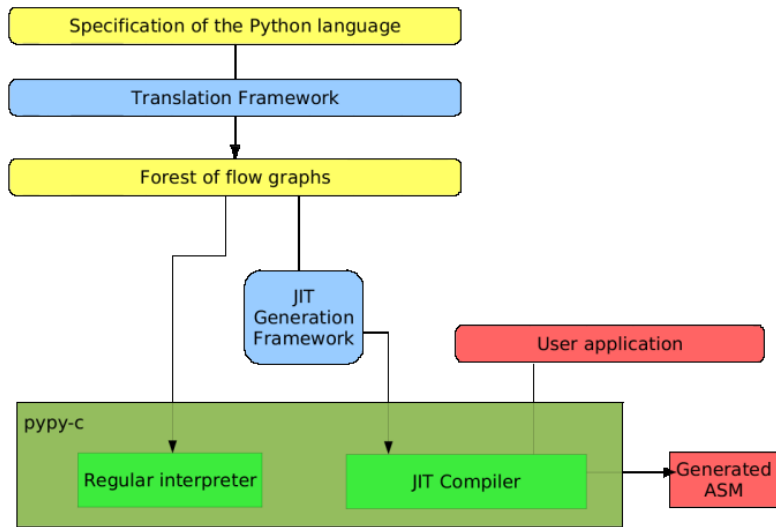
# What is a tracing JIT (history)

- tracing assembler (Dynamo, ~2000)
- tracing Java (~2005)
- tracing JavaScript (~2008)
- PyPy is a “tracing JIT generator”

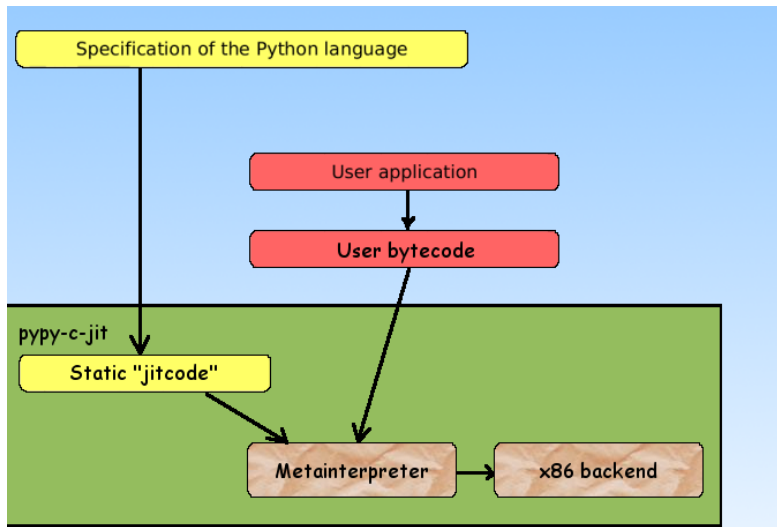
# What is a tracing JIT (history)

- tracing assembler (Dynamo, ~2000)
- tracing Java (~2005)
- tracing JavaScript (~2008)
- PyPy is a “tracing JIT generator”

# The architecture of PyPy



# The architecture of PyPy (2)



# Speed of the PyPy JIT

Python programs that are, or are not, nicely handled by the JIT:

- loops, even across many calls, are nicely handled
- loops with very many taken paths are not
  - ▶ e.g. Python programs that look like interpreters
  - ▶ typical in tracing JITs
- bad support so far for generators and recursion

# The optimizations we get

- != optimizations we wrote :-)
- removed frame handling
  - ▶ local variables are in CPU registers or on the C stack
  - ▶ but `sys._getframe()` works correctly
- “virtuals”: temporary objects are not constructed
  - ▶ `e = a + b + c + d`
  - ▶ and much more complex examples
- attribute and method lookups, etc.

# Example

```
def g(a, b):  
    if a < 5:           # 2  
        return -1  
    return a - b       # 3  
  
def f(x):  
    total = 0  
                                # 1  
    for i in range(x):  
        d = g(i, x)  
        total += d       # 4
```

```
ADD EAX, 1  
CMP EAX, EBX  
JNL <guard 1>  
CMP EAX, 0  
JL <guard 2>  
MOV ECX, EAX  
SUB ECX, EBX  
JO <guard 3>  
ADD EDX, ECX  
JO <guard 4>  
JMP
```

# Practical results

- fast :-)
- so far, x86-32 only
- relatively easy to maintain (or port to x86-64, etc.)
- reminder: works transparently for any Python code
  - ▶ or any language (Prolog JIT :-) at PPDP 2010)
- viable alternative to CPython



# Part 3

cpyext

# cpyext

- CPython extension modules in PyPy
- `pypy-c setup.py build`
- included in PyPy 1.3
- still beta
- 50% of the CPython API is supported
  - ▶ enough for 90% of extension modules

# features

- C API written in Python!
- Testable on top of an interpreted py.py
- Written on top of the object space
- Source compatibility
  - ▶ PyString\_AS\_STRING is actually a function call (instead of a macro)

```
@cpython_api([PyObject], Py_ssize_t, error=-1)
def PyDict_Size(space, w_obj):
    return space.int_w(space.len(w_obj))
```

# implementation

- It was not supposed to work!
  - ▶ different garbage collector
  - ▶ no “borrowed reference”
  - ▶ all the PyObject slots
- not faster than python code!
- PyObject contains ob\_type and ob\_refcnt
  - ▶ The “abstract object interface” is used.
- Some objects contain more:
  - ▶ PyString\_AsString() must keep the buffer alive at a fixed location
  - ▶ PyObject exposes all its fields

# The Reference Counting Issue

- pypy uses a moving garbage collector, starts with static roots to find objects.
- CPython objects don't move, and PyObject\* can point to deallocated memory.
- cpyext builds PyObject as proxies to the “real” interpreter objects
- one dictionary lookup each time the boundary is crossed
- More tricks needed for borrowing references
  - ▶ The object lifetime is tied to its container.
  - ▶ “out of nothing” borrowed references are kept until the end of the current pypy->C call.

# supported modules


- Known to work (after small patches):
  - ▶ wxPython
  - ▶ `_sre`
  - ▶ PyCrypto
  - ▶ PIL
  - ▶ cx\_Oracle
  - ▶ MySQLdb
  - ▶ sqlite

# Why your module will crash

Likely:

```
static PyObject *myException;  
  
void init_foo() {  
    myException = PyException_New(...);  
    Py_AddModule(m, myException); // steals a reference  
}  
  
{  
    PyErr_SetString(myException, "message"); // crash  
}
```

# wxPython on PyPy (1)





# wxPython on PyPy (2)

The screenshot shows the wxPython: (A Demonstration) application window. The left sidebar lists various wxPython controls, with 'I18N' selected. The main area displays an error message: 'An error has occurred while trying to run the demo'. The exception is of type '<type 'exceptions.AttributeError'' with details: 'module' object has no attribute 'getrefcount'. A traceback table follows, showing the sequence of calls from Main.py to I18N.py. At the bottom, a 'Demo Log Messages' panel shows the loading of various demo modules.

wxPython: (A Demonstration)

File Demo Help

wxPython Demos

- wxPython Overview
- Recent Additions/Updates
- RichTextCtrl
- Treebook
- Toolbook
- BitmapFromBuffer
- RawBitmapAccess
- DragScroller
- DelayedResult
- ExpandTextCtrl
- AboutBox
- AlphaDrawing
- GraphicsContext
- CollapsiblePane
- ComboCtrl
- OwnerDrawnComboBox
- BitmapComboBox
- I18N**
- Img2PyArtProvider
- SearchCtrl
- SizedControls
- AUI\_MDI
- TreeMixIn
- AdjustChannels
- RendererNative
- PlateButton
- ResizeWidget
- Cairo
- Cairo\_Snippets
- AdvancedSplash
- AquaButton
- BalloonTip
- ButtonPanel
- CubeColourDialog
- CustomTreeCtrl
- FlatMenu

Filter Demos:

Traceback:

| File... | Line | Function                 | Code  |
|---------|------|--------------------------|---|
| Main.py | 1901 | RunModule                | self.demoPage = module.runTest(self, self.nb, self) |
| I18N.py | 214  | runTest                  | win = LanguageSelectPanel(nb, log)                  |
| I18N.py | 165  | __init__                 | self.OnLangSelectAndTranslate()                     |
| I18N.py | 202  | OnLangSelectAndTranslate | self.updateLanguage(lang)                           |
| I18N.py | 175  | updateLanguage           | assert sys.getrefcount(self.locale) <= 2            |

Entries from the demo module are shown in blue  
Double-click on them to go to the offending line

Demo Log Messages

- OnActivate: True
- OnActivate: False
- OnActivate: True
- Loading demo RendererNative.py...
- Running demo module...
- Loading demo AdjustChannels.py...
- Running demo module...
- Loading demo TreeMixIn.py...
- Loading demo AUI\_MDI.py...

# performance (1)

```
from cx_Oracle import connect, STRING
c = connect('scott/tiger@db')
cur = c.cursor()
var = cur.var(STRING)
```

```
def f():
    for i in range(10000):
        var.setvalue(0, str(i))
        var.getvalue(0)
```

```
python -m timeit -s "from test_oracle import f" "f()"
python2.6:      8.25 msec per loop
pypy-c:        161      msec per loop
pypy-c-jit:    121      msec per loop
```

## performance (2)

Compare with:

```
def f():  
    for i in range(10000):  
        x = str(i)  
        y = int(x)
```

python2.6: 8.18 msec per loop

pypy-c-jit: 1.22 msec per loop

# Future developments

- Some care about speed
  - ▶ The JIT can help to remove (some of) the overhead
- Fill missing API functions (when needed)
- Better support of the PyTypeObject slots
- Think about threads and the GIL
- Think about reference cycles

# Contact / Q&A

- Antonio Cuni: at <http://merlinux.eu>
- Armin Rigo: [arigo \(at\) tunes.org](mailto:arigo@tunes.org)
- Amaury Forgeot d'Arc: [amauryfa \(at\) gmail](mailto:amauryfa@gmail.com)
- And the `#pypy` IRC channel on [freenode.net](http://freenode.net)!
- Links:
  - ▶ PyPy: <http://pypy.org/>
  - ▶ PyPy speed center: <http://speed.pypy.org/>
  - ▶ Blog: <http://morepypy.blogspot.com>