

The Story of Stackless Python

Christian Tismer, Hervé Coatanhay

EuroPython 2012

July 4 2012



About This Talk

- first talk after a long break
 - ▶ *rst2beamer* for the first time

guest speaker:

- Herve Coatanhay about Nagare
 - ▶ PowerPoint (Mac)

Meanwhile I used

- Powerpoint (PC)
- Keynote (Mac)
- Google Docs

poll: What is your favorite slide tool?

About This Talk

- first talk after a long break
 - ▶ *rst2beamer* for the first time

guest speaker:

- Herve Coatanhay about Nagare
 - ▶ PowerPoint (Mac)

Meanwhile I used

- Powerpoint (PC)
- Keynote (Mac)
- Google Docs

poll: What is your favorite slide tool?

About This Talk

- first talk after a long break
 - *rst2beamer* for the first time

guest speaker:

- Herve Coatanhay about Nagare
 - PowerPoint (Mac)

Meanwhile I used

- Powerpoint (PC)
- Keynote (Mac)
- Google Docs

poll: What is your favorite slide tool?

What is Stackless?

- *Stackless is a Python version that does not use the C stack*
 - ▶ really? naah
- Stackless is a Python version that does not keep state on the C stack
 - ▶ the stack *is* used but
 - ▶ cleared between function calls
- Remark:
 - ▶ theoretically. In practice...
 - ▶ ... it is reasonable 90 % of the time
 - ▶ we come back to this!

What is Stackless?

- *Stackless is a Python version that does not use the C stack*
 - ▶ really? naah
- Stackless is a Python version that does not keep state on the C stack
 - ▶ the stack *is* used but
 - ▶ cleared between function calls
- Remark:
 - ▶ theoretically. In practice...
 - ▶ ... it is reasonable 90 % of the time
 - ▶ we come back to this!

What is Stackless?

- *Stackless is a Python version that does not use the C stack*
 - ▶ really? naah
- Stackless is a Python version that does not keep state on the C stack
 - ▶ the stack *is* used but
 - ▶ cleared between function calls
- Remark:
 - ▶ theoretically. In practice...
 - ▶ ... it is reasonable 90 % of the time
 - ▶ we come back to this!

What is Stackless?

- *Stackless is a Python version that does not use the C stack*
 - ▶ really? naah
- Stackless is a Python version that does not keep state on the C stack
 - ▶ the stack *is* used but
 - ▶ cleared between function calls
- Remark:
 - ▶ theoretically. In practice...
 - ▶ ... it is reasonable 90 % of the time
 - ▶ we come back to this!

What is Stackless about?

- it is like CPython
- it can do a little bit more
- adds a single builtin module

```
import stackless
```

- is like an extension
 - ▶ but, sadly, not really
 - ▶ stackless **must** be builtin
 - ▶ **but:** there is a solution...

What is Stackless about?

- it is like CPython
- it can do a little bit more
- adds a single builtin module

```
import stackless
```

- is like an extension
 - ▶ but, sadly, not really
 - ▶ stackless **must** be builtin
 - ▶ **but:** there is a solution...

What is Stackless about?

- it is like CPython
- it can do a little bit more
- adds a single builtin module

```
import stackless
```

- is like an extension
 - ▶ but, sadly, not really
 - ▶ stackless **must** be builtin
 - ▶ **but:** there is a solution...

What is Stackless about?

- it is like CPython
- it can do a little bit more
- adds a single builtin module

```
import stackless
```

- is like an extension
 - ▶ but, sadly, not really
 - ▶ stackless **must** be builtin
 - ▶ **but:** there is a solution...

What is Stackless about?

- it is like CPython
- it can do a little bit more
- adds a single builtin module

```
import stackless
```

- is like an extension
 - ▶ but, sadly, not really
 - ▶ stackless **must** be builtin
 - ▶ **but:** there is a solution...

Now, what is it really about?

- have tiny little “main” programs
 - ▶ `tasklet`
- tasklets communicate via messages
 - ▶ `channel`
- tasklets are often called `microthreads`
 - ▶ but there are no threads at all
 - ▶ only one tasklets runs at any time
- *but see the PyPy STM approach*
 - ▶ this will apply to tasklets as well

Now, what is it really about?

- have tiny little “main” programs
 - ▶ `tasklet`
- tasklets communicate via messages
 - ▶ `channel`
- tasklets are often called `microthreads`
 - ▶ but there are no threads at all
 - ▶ only one tasklets runs at any time
- *but see the PyPy STM approach*
 - ▶ this will apply to tasklets as well

Now, what is it really about?

- have tiny little “main” programs
 - ▶ `tasklet`
- tasklets communicate via messages
 - ▶ `channel`
- tasklets are often called `microthreads`
 - ▶ but there are no threads at all
 - ▶ only one tasklets runs at any time
- *but see the PyPy STM approach*
 - ▶ this will apply to tasklets as well

Now, what is it really about?

- have tiny little “main” programs
 - ▶ `tasklet`
- tasklets communicate via messages
 - ▶ `channel`
- tasklets are often called `microthreads`
 - ▶ but there are no threads at all
 - ▶ only one tasklets runs at any time
- *but see the PyPy STM approach*
 - ▶ this will apply to tasklets as well

Cooperative Multitasking ...

```
>>> import stackless
>>>
>>> channel = stackless.channel()

>>> def receiving_tasklet():
...     print "Receiving tasklet started"
...     print channel.receive()
...     print "Receiving tasklet finished"

>>> def sending_tasklet():
...     print "Sending tasklet started"
...     channel.send("send from sending_tasklet")
...     print "sending tasklet finished"
```

Cooperative Multitasking ...

```
>>> import stackless
>>>
>>> channel = stackless.channel()

>>> def receiving_tasklet():
...     print "Receiving tasklet started"
...     print channel.receive()
...     print "Receiving tasklet finished"

>>> def sending_tasklet():
...     print "Sending tasklet started"
...     channel.send("send from sending_tasklet")
...     print "sending tasklet finished"
```

Cooperative Multitasking ...

```
>>> import stackless
>>>
>>> channel = stackless.channel()

>>> def receiving_tasklet():
...     print "Receiving tasklet started"
...     print channel.receive()
...     print "Receiving tasklet finished"

>>> def sending_tasklet():
...     print "Sending tasklet started"
...     channel.send("send from sending_tasklet")
...     print "sending tasklet finished"
```

... Cooperative Multitasking ...

```
>>> def another_tasklet():  
...     print "Just another tasklet in the scheduler"  
  
>>> stackless.tasklet(receiving_tasklet)()  
<stackless.tasklet object at 0x00A45B30>  
>>> stackless.tasklet(sending_tasklet)()  
<stackless.tasklet object at 0x00A45B70>  
>>> stackless.tasklet(another_tasklet)()  
<stackless.tasklet object at 0x00A45BF0>
```

... Cooperative Multitasking ...

```
>>> def another_tasklet():  
...     print "Just another tasklet in the scheduler"  
  
>>> stackless.tasklet(receiving_tasklet)()  
<stackless.tasklet object at 0x00A45B30>  
>>> stackless.tasklet(sending_tasklet)()  
<stackless.tasklet object at 0x00A45B70>  
>>> stackless.tasklet(another_tasklet)()  
<stackless.tasklet object at 0x00A45BF0>
```

... Cooperative Multitasking

```
<stackless.tasklet object at 0x00A45B70>
>>> stackless.tasklet(another_tasklet)()
<stackless.tasklet object at 0x00A45BF0>
>>>
>>> stackless.run()
Receiving tasklet started
Sending tasklet started
send from sending_tasklet
Receiving tasklet finished
Just another tasklet in the scheduler
sending tasklet finished
```

Why not just the greenlet ?

- greenlets are a subset of stackless
 - ▶ can partially emulate stackless
 - ▶ there is no builtin scheduler
 - ▶ technology quite close to Stackless 2.0
- greenlets are about 10x slower to switch context because using only hard-switching
 - ▶ but that's ok in most cases
- greenlets are kind-of perfect
 - ▶ near zero maintenace
 - ▶ minimal interface
- but the main difference is ...

Why not just the greenlet ?

- greenlets are a subset of stackless
 - ▶ can partially emulate stackless
 - ▶ there is no builtin scheduler
 - ▶ technology quite close to Stackless 2.0
- greenlets are about 10x slower to switch context because using only hard-switching
 - ▶ but that's ok in most cases
- greenlets are kind-of perfect
 - ▶ near zero maintenace
 - ▶ minimal interface
- but the main difference is ...

Why not just the greenlet ?

- greenlets are a subset of stackless
 - ▶ can partially emulate stackless
 - ▶ there is no builtin scheduler
 - ▶ technology quite close to Stackless 2.0
- greenlets are about 10x slower to switch context because using only hard-switching
 - ▶ but that's ok in most cases
- greenlets are kind-of perfect
 - ▶ near zero maintenace
 - ▶ minimal interface
- but the main difference is ...

Why not just the greenlet ?

- greenlets are a subset of stackless
 - ▶ can partially emulate stackless
 - ▶ there is no builtin scheduler
 - ▶ technology quite close to Stackless 2.0
- greenlets are about 10x slower to switch context because using only hard-switching
 - ▶ but that's ok in most cases
- greenlets are kind-of perfect
 - ▶ near zero maintenace
 - ▶ minimal interface
- but the main difference is ...

Excurs: Hard-Switching

Sorry ;-)

Switching program state “the hard way”:

Without notice of the interpreter

- the machine stack gets hijacked
 - ▶ Brute-Force: replace the stack with another one
 - ▶ like threads
- stackless, greenlets
 - ▶ stack slicing
 - ▶ semantically same effect
- switching works fine
- pickling does not work, opaque data on the stack
 - ▶ this is more sophisticated in PyPy, another story...

Excurs: Soft-Switching

Switching program state “the soft way”:

With knowledge of the interpreter

- most efficient implementation in Stackless 3.1
- demands the most effort of the developers
- no opaque data on the stack, pickling does work
 - ▶ again, this is more sophisticated in PyPy
- now we are at the main difference, as you guessed ...

Excurs: Soft-Switching

Switching program state “the soft way”:

With knowledge of the interpreter

- most efficient implementation in Stackless 3.1
- demands the most effort of the developers
- no opaque data on the stack, pickling does work
 - ▶ again, this is more sophisticated in PyPy
- now we are at the main difference, as you guessed ...

Pickling Program State

Persistence (p. 1 of 2)

```
import pickle, sys
import stackless

ch = stackless.channel()

def recurs(depth, level=1):
    print 'enter level %s%d' % (level*' ', level)
    if level >= depth:
        ch.send('hi')
    if level < depth:
        recurs(depth, level+1)
    print 'leave level %s%d' % (level*' ', level)
```

remember to show it interactively

Pickling Program State

Persistence (p. 2 of 2)

```
def demo(depth):  
    t = stackless.tasklet(recurs)(depth)  
    print ch.receive()  
    pickle.dump(t, file('tasklet.pickle', 'wb'))  
  
if __name__ == '__main__':  
    if len(sys.argv) > 1:  
        t = pickle.load(file(sys.argv[1], 'rb'))  
        t.insert()  
    else:  
        t = stackless.tasklet(demo)(9)  
        stackless.run()
```

remember to show it interactively

Script Output 1

```
$ ~/src/stackless/python.exe demo/pickledtasklet.py
enter level 1
enter level 2
enter level 3
enter level 4
enter level 5
enter level 6
enter level 7
enter level 8
enter level 9
hi
leave level 9
leave level 8
leave level 7
leave level 6
leave level 5
leave level 4
leave level 3
leave level 2
leave level 1
```

Script Output 2

```
$ ~/src/stackless/python.exe demo/pickledtasklet.py tasklet.pickle
leave level          9
leave level          8
leave level          7
leave level          6
leave level          5
leave level          4
leave level          3
leave level          2
leave level          1
```

Greenlet vs. Stackless

- Greenlet is a pure extension module
 - ▶ but performance is good enough
- Stackless can pickle program state
 - ▶ but stays a replacement of Python
- Greenlet never can, as an extension
- *easy installation* lets people select greenlet over stackless
 - ▶ see for example the *eventlet* project
 - ▶ *but there is a simple work-around, we'll come to it*
- *they both have their application domains and they will persist.*

Greenlet vs. Stackless

- Greenlet is a pure extension module
 - ▶ but performance is good enough
- Stackless can pickle program state
 - ▶ but stays a replacement of Python
- Greenlet never can, as an extension
- *easy installation* lets people select greenlet over stackless
 - ▶ see for example the *eventlet* project
 - ▶ *but there is a simple work-around, we'll come to it*
- *they both have their application domains and they will persist.*

Greenlet vs. Stackless

- Greenlet is a pure extension module
 - ▶ but performance is good enough
- Stackless can pickle program state
 - ▶ but stays a replacement of Python
- Greenlet never can, as an extension
- *easy installation* lets people select greenlet over stackless
 - ▶ see for example the *eventlet* project
 - ▶ *but there is a simple work-around, we'll come to it*
- *they both have their application domains and they will persist.*

Greenlet vs. Stackless

- Greenlet is a pure extension module
 - ▶ but performance is good enough
- Stackless can pickle program state
 - ▶ but stays a replacement of Python
- Greenlet never can, as an extension
- *easy installation* lets people select greenlet over stackless
 - ▶ see for example the *eventlet* project
 - ▶ *but there is a simple work-around, we'll come to it*
- *they both have their application domains and they will persist.*

Greenlet vs. Stackless

- Greenlet is a pure extension module
 - ▶ but performance is good enough
- Stackless can pickle program state
 - ▶ but stays a replacement of Python
- Greenlet never can, as an extension
- *easy installation* lets people select greenlet over stackless
 - ▶ see for example the *eventlet* project
 - ▶ *but there is a simple work-around, we'll come to it*
- *they both have their application domains and they will persist.*

Why Stackless makes a Difference

- Microthreads ?

- ▶ the feature where I put most effort into
- ▶ can be emulated: (in decreasing speed order)
 - ★ generators (incomplete, “half-sided”)
 - ★ greenlet
 - ★ threads (even ;-)

- Pickling program state ! ==

- persistence

Why Stackless makes a Difference

- Microthreads ?

- ▶ the feature where I put most effort into
- ▶ can be emulated: (in decreasing speed order)
 - ★ generators (incomplete, “half-sided”)
 - ★ greenlet
 - ★ threads (even ;-)

- Pickling program state ! ==

- persistence

Why Stackless makes a Difference

- Microthreads ?

- ▶ the feature where I put most effort into
- ▶ can be emulated: (in decreasing speed order)
 - ★ generators (incomplete, “half-sided”)
 - ★ greenlet
 - ★ threads (even ;-)

- Pickling program state ! ==

- persistence

Why Stackless makes a Difference

- Microthreads ?
 - ▶ the feature where I put most effort into
 - ▶ can be emulated: (in decreasing speed order)
 - ★ generators (incomplete, “half-sided”)
 - ★ greenlet
 - ★ threads (even ;-)
- Pickling program state ! ==
- **persistence**

Persistence, Cloud Computing

- freeze your running program
- let it continue anywhere else
 - ▶ on a different computer
 - ▶ on a different operating system (!)
 - ▶ in a cloud
- migrate your running program
- save snapshots, have checkpoints
 - ▶ without doing any extra-work

Software archeology

- Around since 1998

- ▶ version 1

- ★ using only soft-switching
 - ★ continuation-based
 - ★ *please let me skip old design errors :-)*

- Complete redesign in 2002

- ▶ version 2

- ★ using only hard-switching
 - ★ birth of tasklets and channels

- Concept merge in 2004

- ▶ version 3

- ★ **80-20** rule:
 - ★ soft-switching whenever possible
 - ★ hard-switching if foreign code is on the stack

- ▶ these 80 % can be *pickled* (90?)

- This stayed as version 3.1

Software archeology

- Around since 1998
 - ▶ version 1
 - ★ using only soft-switching
 - ★ continuation-based
 - ★ *please let me skip old design errors :-)*
- Complete redesign in 2002
 - ▶ version 2
 - ★ using only hard-switching
 - ★ birth of tasklets and channels
- Concept merge in 2004
 - ▶ version 3
 - ★ **80-20** rule:
 - ★ soft-switching whenever possible
 - ★ hard-switching if foreign code is on the stack
 - ▶ these 80 % can be *pickled* (90?)
- This stayed as version 3.1

Software archeology

- Around since 1998
 - ▶ version 1
 - ★ using only soft-switching
 - ★ continuation-based
 - ★ *please let me skip old design errors :-)*
- Complete redesign in 2002
 - ▶ version 2
 - ★ using only hard-switching
 - ★ birth of tasklets and channels
- Concept merge in 2004
 - ▶ version 3
 - ★ **80-20** rule:
 - ★ soft-switching whenever possible
 - ★ hard-switching if foreign code is on the stack
 - ▶ these 80 % can be *pickled* (90?)
- This stayed as version 3.1

Status of Stackless Python

- mature
- Python 2 and Python 3, all versions
- maintained by
 - ▶ Richard Tew
 - ▶ Kristjan Valur Jonsson
 - ▶ me (a bit)

The New Direction for Stackless

- `pip install stackless-python`
 - ▶ will install `slpython`
 - ▶ or even `python` (opinions?)
- drop-in replacement of CPython (*psssst*)
- `pip uninstall stackless-python`
 - ▶ Stackless is a bit cheating, as it replaces the `python` binary
 - ▶ but the user perception will be perfect
- *trying stackless made easy!*

The New Direction for Stackless

- `pip install stackless-python`
 - ▶ will install `slpython`
 - ▶ or even `python` (opinions?)
- drop-in replacement of CPython (*psssst*)
- `pip uninstall stackless-python`
 - ▶ Stackless is a bit cheating, as it replaces the `python` binary
 - ▶ but the user perception will be perfect
- *trying stackless made easy!*

The New Direction for Stackless

- `pip install stackless-python`
 - ▶ will install `slpython`
 - ▶ or even `python` (opinions?)
- drop-in replacement of CPython (*psssst*)
- `pip uninstall stackless-python`
 - ▶ Stackless is a bit cheating, as it replaces the `python` binary
 - ▶ but the user perception will be perfect
- *trying stackless made easy!*

New Direction (cont'd)

- first prototype yesterday from Anselm Kruis (*applause*)
 - ▶ works on Windows
 - ▶ OS X
 - ★ I'll do that one
 - ▶ Linux
 - ★ soon as well
- being very careful to stay compatible
 - ▶ python 2.7.3 installs stackless for 2.7.3
 - ▶ python 3.2.3 installs stackless for 3.2.3
 - ▶ python 2.7.2 : *please upgrade* - or maybe have an over-ride option?

New Direction (cont'd)

- first prototype yesterday from Anselm Kruis (*applause*)
 - ▶ works on Windows
 - ▶ OS X
 - ★ I'll do that one
 - ▶ Linux
 - ★ soon as well
- being very careful to stay compatible
 - ▶ python 2.7.3 installs stackless for 2.7.3
 - ▶ python 3.2.3 installs stackless for 3.2.3
 - ▶ python 2.7.2 : *please upgrade* - or maybe have an over-ride option?

New Direction (cont'd)

- first prototype yesterday from Anselm Kruis (*applause*)
 - ▶ works on Windows
 - ▶ OS X
 - ★ I'll do that one
 - ▶ Linux
 - ★ soon as well
- being very careful to stay compatible
 - ▶ python 2.7.3 installs stackless for 2.7.3
 - ▶ python 3.2.3 installs stackless for 3.2.3
 - ▶ python 2.7.2 : *please upgrade* - or maybe have an over-ride option?

New Direction (cont'd)

- first prototype yesterday from Anselm Kruis (*applause*)
 - ▶ works on Windows
 - ▶ OS X
 - ★ I'll do that one
 - ▶ Linux
 - ★ soon as well
- being very careful to stay compatible
 - ▶ python 2.7.3 installs stackless for 2.7.3
 - ▶ python 3.2.3 installs stackless for 3.2.3
 - ▶ python 2.7.2 : *please upgrade* - or maybe have an over-ride option?

Consequences of the Pseudo-Package

The technical effect is almost nothing.

The psychological impact is probably huge:

- stackless is easy to install and uninstall
- people can simply try if it fits their needs
- the never ending discussion
 - ▶ “Why is Stackless not included in the Python core?”
- **has ended**
 - ▶ “Why should we, after all?”
 - ▶ hey Guido :-)
 - ▶ what a relief, for you and me

Consequences of the Pseudo-Package

The technical effect is almost nothing.

The psychological impact is probably huge:

- stackless is easy to install and uninstall
- people can simply try if it fits their needs
- the never ending discussion
 - ▶ “Why is Stackless not included in the Python core?”
- **has ended**
 - ▶ “Why should we, after all?”
 - ▶ hey Guido :-)
 - ▶ what a relief, for you and me

Consequences of the Pseudo-Package

The technical effect is almost nothing.

The psychological impact is probably huge:

- stackless is easy to install and uninstall
- people can simply try if it fits their needs
- the never ending discussion
 - ▶ “Why is Stackless not included in the Python core?”
- **has ended**
 - ▶ “Why should we, after all?”
 - ▶ hey Guido :-)
 - ▶ what a relief, for you and me

Consequences of the Pseudo-Package

The technical effect is almost nothing.

The psychological impact is probably huge:

- stackless is easy to install and uninstall
- people can simply try if it fits their needs
- the never ending discussion
 - ▶ “Why is Stackless not included in the Python core?”
- **has ended**
 - ▶ “Why should we, after all?”
 - ▶ hey Guido :-)
 - ▶ what a relief, for you and me

Consequences of the Pseudo-Package

The technical effect is almost nothing.

The psychological impact is probably huge:

- stackless is easy to install and uninstall
- people can simply try if it fits their needs
- the never ending discussion
 - ▶ “Why is Stackless not included in the Python core?”
- **has ended**
 - ▶ “Why should we, after all?”
 - ▶ hey Guido :-)
 - ▶ what a relief, for you and me

Consequences of the Pseudo-Package

The technical effect is almost nothing.

The psychological impact is probably huge:

- stackless is easy to install and uninstall
- people can simply try if it fits their needs
- the never ending discussion
 - ▶ “Why is Stackless not included in the Python core?”
- **has ended**
 - ▶ “Why should we, after all?”
 - ▶ hey Guido :-)
 - ▶ what a relief, for you and me

Status of Stackless PyPy

- was completely implemented before the Jit
 - ▶ together with greenlets coroutines
 - ▶ not Jit compatible
- was “too complete” with a 30% performance hit
- new approach is almost ready
 - ▶ with full Jit support
 - ▶ but needs some fixing
 - ▶ this *will* be efficient

Applications using Stackless Python

- The Eve Online MMORPG
<http://www.eveonline.com/>
 - ▶ based their games on Stackless since 1998
- science + computing ag, Anselm Kruis
<https://ep2012.europython.eu/conference/p/anselm-kruis>
- The Nagare Web Framework
<http://www.nagare.org/>
 - ▶ works because of Stackless Pickling
- today's majority: persistence

Thank you

- the new Stackless Website
`http://www.stackless.com/`
 - ▶ a **great** donation from Alain Pourier, *Nagare*
- You can hire me as a consultant
- Questions?