

Python and PyPy performance (not) for dummies

Antonio Cuni and Maciej Fijałkowski

EuroPython 2015

July 21, 2015

About us

- PyPy core devs
- `vmprof`, `cffi`, `pdb++`, `fancycompleter`, ...
- Consultants
- <http://baroquesoftware.com/>

Optimization for dummies

- Obligatory citation
 - ▶ *premature optimization is the root of all evil* (D. Knuth)
- Pareto principle, or 80-20 rule
 - ▶ 80% of the time will be spent in 20% of the program
 - ▶ 20% of 1 mln is 200 000
- Two golden rules:
 1. Identify the slow spots
 2. Optimize them

This talk

- Two parts
 1. How to identify the slow spots
 2. How to address the problems

Part 1

- identifying the slow spots

What is performance?

- something quantifiable by numbers
- usually, time spent doing task X
- number of requests, latency, etc.
- statistical properties about that metric

Do you have a performance problem?

- what you're trying to measure
- means to measure it (production, benchmarks, etc.)
- is Python is the cause here?
- environment to quickly measure and check the results
 - ▶ same as for debugging

When Python is the problem

- tools, timers etc.
- systems are too complicated to **guess** which will be faster
- find your bottlenecks
- 20/80 (but 20% of million lines is 200 000 lines, remember that)

Profilers landscape

- cProfile, runSnakeRun (high overhead) - event based profiler
- plop, vmprof - statistical profilers
- cProfile & vmprof work on pypy

vmprof

- inspired by `gperftools`
- statistical profiler run by an interrupt (~300Hz on modern linux)
- sampling the C stack
- CPython, PyPy, possibly more virtual machines

why not gperftools?

- C stack does not contain python-level frames
- 90% `PyEval_EvalFrame` and other internals
- we want python-level functions
- picture is even more confusing in the presence of the JIT

using vmprof

- demo
- `http://vmprof.readthedocs.org`

using vmprof in production

- low overhead (5-10%), possibly lower in the future
- possibility of realtime monitoring (coming)

vmprof future

- profiler as a service
- realtime advanced visualization

Part 2

Make it fast

Tools


- Endless list of tools/techniques to increment speed
- C extension
- Cython
- numba
- "performance tricks"
- **PyPy**
 - ▶ We'll concentrate on it
 - ▶ WARNING: we wrote it, we are biased :)
 - ▶ gives you most wins for free (*)

What is PyPy

- Alternative, fast Python implementation
- Performance: JIT compiler, advanced GC
- PyPy 2.6.0 (Python version 2.7.9)
- Py3k as usual in progress (3.2.5 out, 3.3 in development)
- `http://pypy.org`
- EP Talks:
 - ▶ The GIL is dead: PyPy-STM (July 23, 16:45 by Armin Rigo)
 - ▶ PyPy ecosystem: CFFI, numpy, scipy, etc (July 24, 15:15 by Romain Guillebert)

Speed: 7x faster than CPython

How fast is PyPy?




The JIT

```
def main():  
    init()  
    some_quick_code()  
    for x in large_list:  
        do_something(x)  
    some_other_code()  
    while condition():  
        expensive_computation()
```

The JIT

```
def main():  
    init()  
    some_quick_code()  
    for x in large_list:  
        do_something(x)  
    some_other_code()  
    while condition():  
        expensive_computation()
```


NO JIT




The JIT

```
def main():  
    init()  
    some_quick_code()  
    for x in large_list:  
        do_something(x)  
    some_other_code()  
    while condition():  
        expensive_computation()
```

NO JIT



JIT



JIT overview

- Tracing JIT
 - ▶ detect and compile "hot" code
- **Specialization**
- Precompute as much as possible
- Constant propagation
- Aggressive inlining

Specialization (1)

- `obj.foo()`
- which code is executed? (SIMPLIFIED)
 - ▶ lookup `foo` in `obj.__dict__`
 - ▶ lookup `foo` in `obj.__class__`
 - ▶ lookup `foo` in `obj.__bases__[0]`, etc.
 - ▶ finally, execute `foo`
- without JIT, you need to do these steps again and again
- Precompute the lookup?

Specialization (2)

- pretend and assume that `obj.__class__` IS constant
 - ▶ "promotion"
- guard
 - ▶ check our assumption: if it's false, bail out
- now we can directly jump to `foo` code
 - ▶ ...unless `foo` is in `obj.__dict__`: GUARD!
 - ▶ ...unless `foo.__class__.__dict__` changed: GUARD!
- Too many guard failures?
 - ▶ Compile some more assembler!
- guards are cheap
 - ▶ out-of-line guards even more

Specialization (3)

- who decides what to promote/specialize for?
 - ▶ we, the PyPy devs :)
 - ▶ heuristics
- instance attributes are never promoted
- class attributes are promoted by default (with some exceptions)
- module attributes (i.e., globals) as well
- bytecode constants

Specialization trade-offs

- Too much specialization
 - ▶ guards fails often
 - ▶ explosion of assembler
- Not enough specialization
 - ▶ inefficient code

Guidos points



Guido van Rossum

Shared publicly - Sep 10, 2012

Some patterns for fast Python. Know any others?

- Avoid overengineering datastructures. Tuples are better than objects (try `namedtuple` too though). Prefer simple fields over `getter/setter` functions.
- Built-in datatypes are your friends. Use more numbers, strings, tuples, lists, sets, dicts. Also check out the `collections` library, esp. `deque`.
- Be suspicious of function/method calls; creating a stack frame is expensive.
- Don't write Java (or C++, or Javascript, ...) in Python.
- Are you sure it's too slow? Profile before optimizing!
- The universal speed-up is rewriting small bits of code in C. Do this only when all else fails.

Don't do it on PyPy (or at all)

- simple is better than complicated
- avoid string concatenation in the loop
- avoid replacing simple loop with itertools monsters
- "move stuff to C" is (almost) never a good idea
- use `cffi` when calling C
- avoid C extensions using CPython C API
- avoid creating classes at runtime

Example

- `map(operator.attrgetter('x'), list)`

vs

- `[x.x for x in list]`

More about PyPy

- we are going to run a PyPy open space (tomorrow 18:00 @ A4)
- come ask more questions

Q&A?

- Thank you!
- `http://baroquesoftware.com`
- `http://pypy.org`
- `http://vmprof.readthedocs.org`