

Software Transactional Memory "for real"

Armin Rigo

FSCONS 2012

November 10, 2012



Introduction

- This talk is about programming multi- or many-core machines

About myself

- Armin Rigo
- “Language implementation guy”
- PyPy project
 - ▶ Python in Python
 - ▶ includes a Just-in-Time Compiler “Generator” for Python and any other dynamic language

About myself

- Armin Rigo
- “Language implementation guy”
- PyPy project
 - ▶ Python in Python
 - ▶ includes a Just-in-Time Compiler “Generator” for Python and any other dynamic language

Motivation

- A single-core program is getting exponentially slower than a multi-core one
- Using several processes exchanging data
 - ▶ works fine in some cases
 - ▶ but becomes a large mess in others
- Using several threads
 - ▶ this talk!

Motivation

- A single-core program is getting exponentially slower than a multi-core one
- Using several processes exchanging data
 - ▶ works fine in some cases
 - ▶ but becomes a large mess in others
- Using several threads
 - ▶ this talk!

Motivation

- A single-core program is getting exponentially slower than a multi-core one
- Using several processes exchanging data
 - ▶ works fine in some cases
 - ▶ but becomes a large mess in others
- Using several threads
 - ▶ this talk!

Common solution

- Organize your program in multiple threads
- Add synchronization when accessing shared, non-read-only data

Synchronization with locks

- Carefully place locks around every access to shared data
- How do you know if you missed a place?
 - ▶ hard to catch by writing tests
 - ▶ instead you get obscure rare run-time crashes
- Issues when scaling to a large program
 - ▶ order of acquisition
 - ▶ deadlocks

Synchronization with locks

- Carefully place locks around every access to shared data
- How do you know if you missed a place?
 - ▶ hard to catch by writing tests
 - ▶ instead you get obscure rare run-time crashes
- Issues when scaling to a large program
 - ▶ order of acquisition
 - ▶ deadlocks

Synchronization with locks

- Carefully place locks around every access to shared data
- How do you know if you missed a place?
 - ▶ hard to catch by writing tests
 - ▶ instead you get obscure rare run-time crashes
- Issues when scaling to a large program
 - ▶ order of acquisition
 - ▶ deadlocks

Synchronization with TM

- TM = Transactional Memory

Locks

```
mylock.acquire();  
x = list1.pop();  
list2.append(x);  
mylock.release();
```

Transactional Memory

```
atomic {  
    x = list1.pop();  
    list2.append(x);  
}
```

Synchronization with TM

- TM = Transactional Memory

Locks

```
mylock.acquire();  
x = list1.pop();  
list2.append(x);  
mylock.release();
```

Transactional Memory

```
atomic {  
    x = list1.pop();  
    list2.append(x);  
}
```

Locks versus TM

- Locks




- TM



Locks versus TM

- Locks



- TM in case of conflict



Synchronization with TM

- “Optimistic” approach:
 - ▶ no lock to protect shared data in memory
 - ▶ instead, track all memory accesses
 - ▶ detect actual conflicts
 - ▶ if conflict, restart the whole “transaction”
- Easier to use
 - ▶ no need to name locks
 - ▶ no deadlocks
 - ▶ “composability”

Synchronization with TM

- “Optimistic” approach:
 - ▶ no lock to protect shared data in memory
 - ▶ instead, track all memory accesses
 - ▶ detect actual conflicts
 - ▶ if conflict, restart the whole “transaction”
- Easier to use
 - ▶ no need to name locks
 - ▶ no deadlocks
 - ▶ “composability”

HTM versus STM

- HTM = Hardware Transactional Memory
 - ▶ Intel Haswell CPU, 2013
 - ▶ and others
- STM = Software Transactional Memory
 - ▶ various approaches
 - ▶ large overhead (2x-10x), but getting faster
 - ▶ experimental in PyPy: read/write barriers, as with GC

HTM versus STM

- HTM = Hardware Transactional Memory
 - ▶ Intel Haswell CPU, 2013
 - ▶ and others
- STM = Software Transactional Memory
 - ▶ various approaches
 - ▶ large overhead (2x-10x), but getting faster
 - ▶ experimental in PyPy: read/write barriers, as with GC

The catch

- You Still Have To Use Threads
- Threads are hard to debug, non-reproducible
- Threads are Messy

The catch

- You Still Have To Use Threads
- Threads are hard to debug, non-reproducible
- Threads are Messy

The catch

- You Still Have To Use Threads
- Threads are hard to debug, non-reproducible
- Threads are Messy

The catch

- You Still Have To Use Threads
- Threads are hard to debug, non-reproducible
- Threads are Messy

Issue with threads

- TM does not solve this problem:
- How do you know if you missed a place to put `atomic` around?
 - ▶ hard to catch by writing tests
 - ▶ instead you get obscure rare run-time crashes
- What if we put `atomic` everywhere?

Issue with threads

- TM does not solve this problem:
- How do you know if you missed a place to put `atomic` around?
 - ▶ hard to catch by writing tests
 - ▶ instead you get obscure rare run-time crashes
- What if we put `atomic` everywhere?

Analogy with Garbage Collection

- Explicit Memory Management:
 - ▶ messy, hard to debug rare leaks or corruptions
- Automatic GC solves it
 - ▶ common languages either have a GC or not
 - ▶ if they have a GC, it controls almost *all* objects
 - ▶ not just a small part of them

Analogy with Garbage Collection


- Explicit Memory Management:
 - ▶ messy, hard to debug rare leaks or corruptions
- Automatic GC solves it
 - ▶ common languages either have a GC or not
 - ▶ if they have a GC, it controls almost *all* objects
 - ▶ not just a small part of them

Proposed solution

- Put `atomic` everywhere...
- in other words, Run Everything with TM

Proposed solution

- Really needs TM. With locks, you'd get this:



- With TM you can get this:



In a few words

- Longer transactions
- Corresponding to larger parts of the program
- The underlying multi-threaded model becomes implicit

Typical example

- You want to run $f1()$ and $f2()$ and $f3()$
- Assume they are “mostly independent”
 - ▶ i.e. we expect that we can run them in parallel
 - ▶ but we cannot prove it, we just hope that in the common case we can
- In case of conflicts, we don't want random behavior
 - ▶ i.e. we don't want thread-like non-determinism and crashes

Typical example

- You want to run $f1()$ and $f2()$ and $f3()$
- Assume they are “mostly independent”
 - ▶ i.e. we expect that we can run them in parallel
 - ▶ but we cannot prove it, we just hope that in the common case we can
- In case of conflicts, we don't want random behavior
 - ▶ i.e. we don't want thread-like non-determinism and crashes

Typical example

- You want to run $f1()$ and $f2()$ and $f3()$
- Assume they are “mostly independent”
 - ▶ i.e. we expect that we can run them in parallel
 - ▶ but we cannot prove it, we just hope that in the common case we can
- In case of conflicts, we don't want random behavior
 - ▶ i.e. we don't want thread-like non-determinism and crashes

Pooling and atomic statements

- Solution: use a library that creates a pool of threads
- Each thread picks a function from the list and runs it with `atomic`

Results

- The behavior is “as if” we had run $f1()$, $f2()$ and $f3()$ sequentially
- The programmer chooses if he wants this fixed order, or if any order is fine
- Threads are hidden from the programmer

More generally

- This was an example only
- **TM gives various new ways to hide threads under a nice interface**

Not the Ultimate Solution

- Much easier for the programmer to get reproducible results
- But maybe too many conflicts
- “The right side” of the problem
 - ▶ start with a working program, and improve performance
 - ▶ as opposed to: with locks, start with a fast program, and debug crashes
 - ▶ we will need new debugging tools

Not the Ultimate Solution

- Much easier for the programmer to get reproducible results
- But maybe too many conflicts
- “The right side” of the problem
 - ▶ start with a working program, and improve performance
 - ▶ as opposed to: with locks, start with a fast program, and debug crashes
 - ▶ we will need new debugging tools

- PyPy-STM: a version of PyPy with Software Transactional Memory
 - ▶ in-progress, but basically working
 - ▶ solves the “GIL issue” but more importantly adds `atomic`
- `http://pypy.org/`
- Thank you!

- PyPy-STM: a version of PyPy with Software Transactional Memory
 - ▶ in-progress, but basically working
 - ▶ solves the “GIL issue” but more importantly adds `atomic`
- `http://pypy.org/`
- Thank you!