

Faster than C#: efficient implementation of dynamic languages on .NET

Antonio Cuni
Davide Ancona
Armin Rigo

ICOOOLPS@ECOOP 2009 - Genova, Italy

July 6, 2009

Introduction

- Dynamic languages are nice
 - ▶ e.g., Python
- so are .NET and the JVM
- Problem: slow!
- Solution: make them faster :-)
- We concentrate our efforts on .NET

State of the art

- IronPython
- Jython
- JRuby, Groovy, ...
- Self
- Javascript: TraceMonkey, V8
- ...

State of the art

- IronPython
- Jython
- JRuby, Groovy, ...
- **Self**
- Javascript: TraceMonkey, V8
- ...

Why so slow?

- Hard to compile efficiently
- Lack of type information at compile-time
- VMs not optimized to run them
- .NET is a multi-language VM?
 - ▶ Sure, as long as the language is C#
- JVM is in a better shape, but still heavily optimized for Java

Why so slow?

- Hard to compile efficiently
- Lack of type information at compile-time
- VMs not optimized to run them
- .NET is a multi-language VM?
 - ▶ Sure, as long as the language is C#
- JVM is in a better shape, but still heavily optimized for Java

Why so slow?

- Hard to compile efficiently
- Lack of type information at compile-time
- VMs not optimized to run them
- .NET is a multi-language VM?
 - ▶ Sure, as long as the language is C#
- JVM is in a better shape, but still heavily optimized for Java

JIT compiler

- Wait until you know what you need
- Interweave compile-time and runtime
- Exploit runtime information

JIT on top of .NET

- JIT layering
- How to extend existing code?
- Fight the VM

JIT compiler

- Wait until you know what you need
- Interweave compile-time and runtime
- Exploit runtime information

JIT on top of .NET

- JIT layering
- How to extend existing code?
- Fight the VM

PyPy

- Python in Python
 - (lots of features and goals)
 - **JIT compiler generator**
 - Python semantics for free
 - JIT frontend
 - ▶ Not limited to Python
 - JIT backends
 - ▶ x86 backend
 - ▶ **CLI/.NET backend**
- Note: this talk is about JIT v2

PyPy

- Python in Python
- (lots of features and goals)
- **JIT compiler generator**
- Python semantics for free
- JIT frontend
 - ▶ Not limited to Python
- JIT backends
 - ▶ x86 backend
 - ▶ **CLI/.NET backend**
- Note: this talk is about JIT v2

Partial evaluation (PE)

- Assume the Python bytecode to be constant
- Constant-propagate it into the Python interpreter.
- Colors
 - ▶ **Green**: compile-time value
 - ▶ **Red**: runtime value

Partial Evaluation with Colors

- **Green operations:** unchanged, executed at compile-time
- **Red operations:** converted into corresponding code emitting code

Example

```
def f(x, y):  
    x2 = x * x  
    y2 = y * y  
    return x2 + y2
```

case x=10

```
def f_10(y):  
    y2 = y * y  
    return 100 + y2
```

Partial Evaluation with Colors

- **Green operations:** unchanged, executed at compile-time
- **Red operations:** converted into corresponding code emitting code

Example

```
def f(x, y):  
    x2 = x * x  
    y2 = y * y  
    return x2 + y2
```

case x=10

```
def f_10(y):  
    y2 = y * y  
    return 100 + y2
```

Partial Evaluation with Colors

- **Green operations:** unchanged, executed at compile-time
- **Red operations:** converted into corresponding code emitting code

Example

```
def f(x, y):  
    x2 = x * x  
    y2 = y * y  
    return x2 + y2
```

case x=10

```
def f_10(y):  
    y2 = y * y  
    return 100 + y2
```

Challenges

- A shortcoming of PE is that in many cases not much can be really assumed constant at compile-time: poor results
- Effective dynamic compilation requires feedback of runtime information into compile-time
- For a dynamic language: types are a primary example

Solution: Promotion

- “Promote” run-time values to compile-time
- Promotion guided by few hints in the interpreter
- Stop the compilation at promotions
- Execute until promotion points
- Compile more

Promotion (example)

Example

```
def f(x, y):  
    x1 = hint(x, promote=True)  
    return x1*x1 + y*y
```

original

```
def f_(x, y):  
    switch x:  
        pass  
    default:  
        compile_more(x)
```

augmented

```
def f_(x, y):  
    switch x:  
        case 3:  
            return 9 + y*y  
    default:  
        compile_more(x)
```

Promotion (example)

Example

```
def f(x, y):  
    x1 = hint(x, promote=True)  
    return x1*x1 + y*y
```

original

```
def f_(x, y):  
    switch x:  
        pass  
    default:  
        compile_more(x)
```

augmented

```
def f_(x, y):  
    switch x:  
        case 3:  
            return 9 + y*y  
    default:  
        compile_more(x)
```

Promotion (example)

Example

```
def f(x, y):  
    x1 = hint(x, promote=True)  
    return x1*x1 + y*y
```

original

```
def f_(x, y):  
    switch x:  
        pass  
    default:  
        compile_more(x)
```


augmented

```
def f_(x, y):  
    switch x:  
        case 3:  
            return 9 + y*y  
    default:  
        compile_more(x)
```



Promotion on .NET

- Flexswitch
 - ▶ Growable switch
 - ▶ Can add new cases at runtime
- Ideally as efficient as a jump
- No support from the VM
- Very costly
- Still effective as long as it's not in the hot path

Flexswitch example



Flexswitch example



Flexswitch for CLI

- Unit of compilation: method
- Flowgraphs split into multiple methods
- Primary method
 - ▶ Contains a trampoline
 - ▶ Array of delegates
- Secondary methods
 - ▶ Stored into that array
- Jumps between secondary methods go through the trampoline
- Hard (and slow!) to pass arguments around

TLC

- Python not (yet) supported :-)
- Dynamic toy language
- Designed to be “as slow as Python”
- Stack manipulation
- Boxed integers
- Dynamic lookup of methods

Benchmarks (1)

	Factorial			
<i>n</i>	10	10^7	10^8	10^9
Interp	0.031	30.984	N/A	N/A
JIT	0.422	0.453	0.859	4.844
JIT 2	0.000	0.047	0.453	4.641
C#	0.000	0.031	0.359	3.438
Interp/JIT 2	N/A	661.000	N/A	N/A
JIT 2/C#	N/A	1.500	1.261	1.350

Benchmarks (2)

	Fibonacci			
<i>n</i>	10	10^7	10^8	10^9
Interp	0.031	29.359	N/A	N/A
JIT	0.453	0.469	0.688	2.953
JIT 2	0.000	0.016	0.250	2.500
C#	0.000	0.016	0.234	2.453
Interp/JIT 2	N/A	1879.962	N/A	N/A
JIT 2/C#	N/A	0.999	1.067	1.019

Benchmars (3)

```
def main(n):  
    if n < 0:  
        n = -n  
        obj = new(value, accumulate=count)  
    else:  
        obj = new(value, accumulate=add)  
    obj.value = 0  
    while n > 0:  
        n = n - 1  
        obj.accumulate(n)  
    return obj.value  
  
def count(x):  
    this.value = this.value + 1  
def add(x):  
    this.value = this.value + x
```

Benchmars (4)

<i>n</i>	Accumulator			
	10	10^7	10^8	10^9
Interp	0.031	43.063	N/A	N/A
JIT	0.453	0.516	0.875	4.188
JIT 2	0.000	0.047	0.453	3.672
C#	0.000	0.063	0.563	5.953
Interp/JIT 2	N/A	918.765	N/A	N/A
JIT 2/C#	N/A	0.750	0.806	0.617

Future work

- Non local jumps are terribly slow
- Good results only if they are not in the inner loop
- Recompile hot non-local jumps?
- Tracing JIT?
 - ▶ You have just seen it in the previous talk :-)

Contributions

- JIT layering works
 - ▶ Optimize different levels of overhead
 - ▶ .NET's own JIT could be improved
- Current VMs are limited
 - ▶ How to make them more friendly to dynamic languages?