

# Runtime Feedback in a Meta-Tracing JIT for Efficient Dynamic Languages

Carl Friedrich Bolz<sup>1</sup>   Antonio Cuni<sup>3</sup>   Maciej Fijałkowski<sup>2</sup>  
Michael Leuschel<sup>1</sup>   Samuele Pedroni<sup>3</sup>   Armin Rigo<sup>1</sup>

<sup>1</sup>Heinrich-Heine-Universität Düsseldorf, STUPS Group, Germany

<sup>2</sup>merlinux GmbH, Hildesheim, Germany

<sup>3</sup>Open End, Göteborg, Sweden

ICOOOLPS 2011, July 26, 2011

# Good JIT Compilers for Dynamic Languages are Hard

- recent languages like Python, Ruby, JS have complex core semantics
- many corner cases, even hard to interpret correctly

# Good JIT Compilers for Dynamic Languages are Hard

- recent languages like Python, Ruby, JS have complex core semantics
- many corner cases, even hard to interpret correctly
- feedback of runtime information to the compiler is necessary
- correct exploitation of this information in the compiler


# Good JIT Compilers for Dynamic Languages are Hard

- recent languages like Python, Ruby, JS have complex core semantics
- many corner cases, even hard to interpret correctly
- feedback of runtime information to the compiler is necessary
- correct exploitation of this information in the compiler

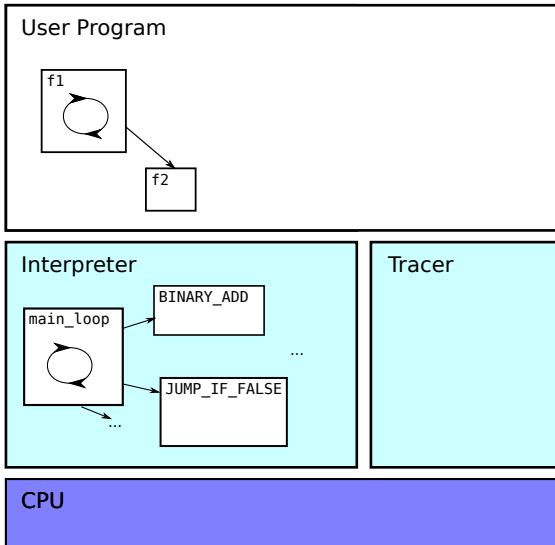
## Problems

- 1 implement all corner-cases of semantics correctly
- 2 ... and the common cases efficiently
- 3 feed back and exploit runtime information

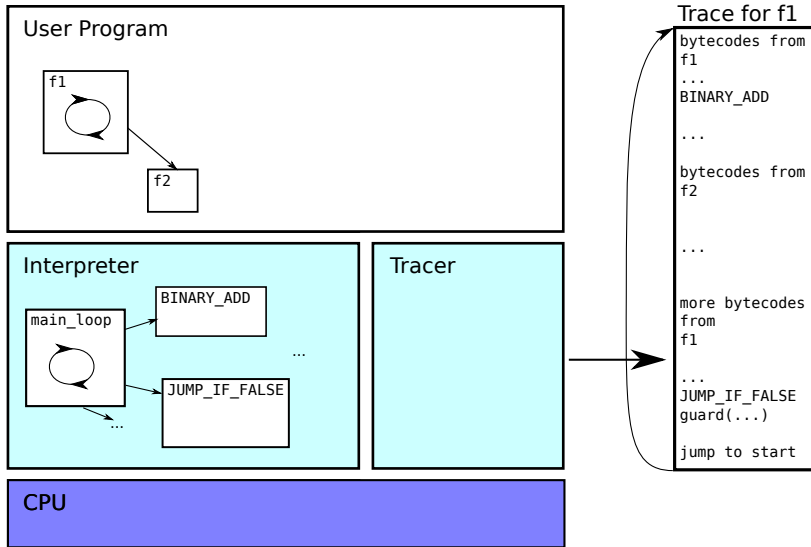
# An Interpreter



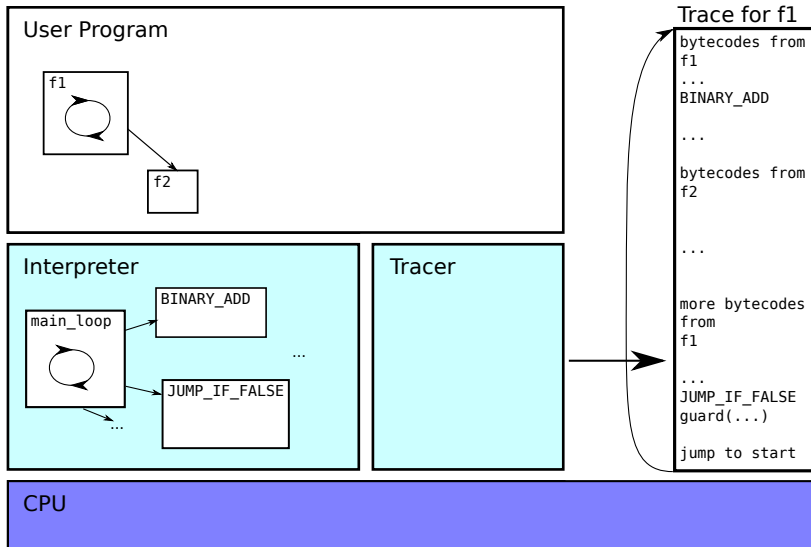
# A Tracing JIT



# A Tracing JIT



# A Tracing JIT





## Advantages:

- can be added to existing VM
- interpreter does a lot of work
- can fall back to interpreter for uncommon paths

# Granularity Problems

- if the tracer records bytecode, not enough information is there
- many dynamic languages have bytecodes that contain complex logic
- need to expand the bytecode in the trace into something more explicit
- this duplicates the language semantics in the tracer/optimizer

# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

- check for `x.__getattr__`, if there, call it

# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

- check for `x.__getattr__`, if there, call it
- look for the attribute in the object's dictionary, if it's there, return it

# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

- check for `x.__getattr__`, if there, call it
- look for the attribute in the object's dictionary, if it's there, return it
- walk up the MRO and look in each class' dictionary for the attribute

# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

- check for `x.__getattr__`, if there, call it
- look for the attribute in the object's dictionary, if it's there, return it
- walk up the MRO and look in each class' dictionary for the attribute
- if the attribute is found, call its `__get__` attribute and return the result

# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

- check for `x.__getattr__`, if there, call it
- look for the attribute in the object's dictionary, if it's there, return it
- walk up the MRO and look in each class' dictionary for the attribute
- if the attribute is found, call its `__get__` attribute and return the result
- if the attribute is not found, look for `x.__getattr__`, if there, call it



# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

- check for `x.__getattr__`, if there, call it
- look for the attribute in the object's dictionary, if it's there, return it
- walk up the MRO and look in each class' dictionary for the attribute
- if the attribute is found, call its `__get__` attribute and return the result
- if the attribute is not found, look for `x.__getattr__`, if there, call it
- raise an `AttributeError`

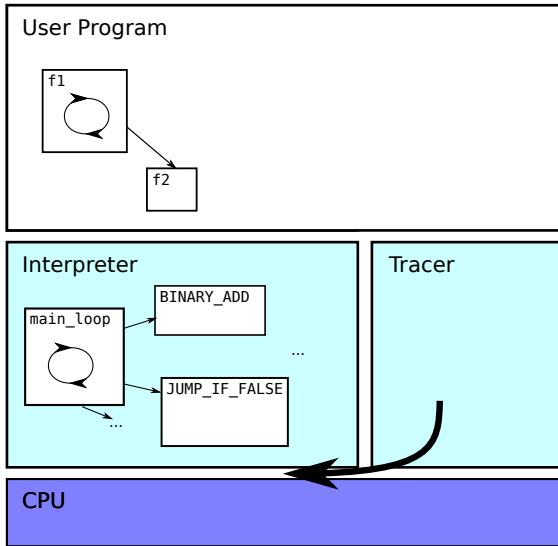
# Example: Attribute Reads in Python

What happens when an attribute `x.m` is read? (simplified)

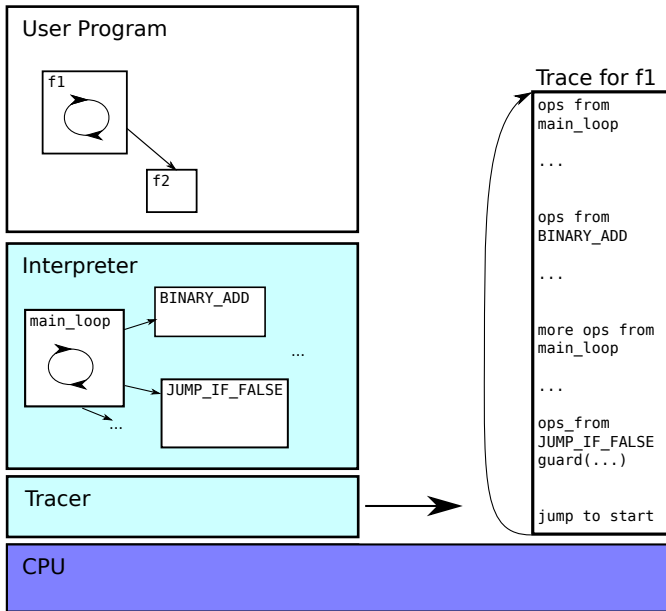
- check for `x.__getattr__`, if there, call it
- look for the attribute in the object's dictionary, if it's there, return it
- walk up the MRO and look in each class' dictionary for the attribute
- if the attribute is found, call its `__get__` attribute and return the result
- if the attribute is not found, look for `x.__getattr__`, if there, call it
- raise an `AttributeError`

all this is one bytecode

# Idea of Meta-Tracing



# Meta-Tracing



## Advantages:

- semantics are always like that of the interpreter
- trace fully contains language semantics
- meta-tracers can be reused for various interpreters

## Advantages:

- semantics are always like that of the interpreter
- trace fully contains language semantics
- meta-tracers can be reused for various interpreters

a few meta-tracing systems have been built:

- Sullivan et.al. describe a meta-tracer using the Dynamo RIO system
- Yermolovich et.al. run a Lua implementation on top of a tracing JS implementation
- SPUR is a tracing JIT for CLR bytecodes, which is used to speed up a JS implementation in C#

A general environment for implementing dynamic languages

A general environment for implementing dynamic languages

## Approach

- write an interpreter for the language in RPython
- compilable to an efficient C-based VM



A general environment for implementing dynamic languages

## Approach

- write an interpreter for the language in RPython
- compilable to an efficient C-based VM
- (RPython is a restricted subset of Python)

A general environment for implementing dynamic languages

## Approach

- write an interpreter for the language in RPython
- compilable to an efficient C-based VM
- (RPython is a restricted subset of Python)

## PyPy's Meta-Tracing JIT

- PyPy contains a meta-tracing JIT for interpreters in RPython
- needs a few source-code hints (or annotations) in the interpreter
- powerful general optimizations

## Problems of Naive Meta-Tracing:

- no runtime feedback of user-level types
- tracer does not know about invariants in the interpreter

## Problems of Naive Meta-Tracing:

- no runtime feedback of user-level types
- tracer does not know about invariants in the interpreter

## Proposed Solutions

- introduce *hints* that the interpreter-author can use
- hints are annotation in the interpreter
- they give information to the meta-tracer

## Problems of Naive Meta-Tracing:

- no runtime feedback of user-level types
- tracer does not know about invariants in the interpreter


## Proposed Solutions

- introduce *hints* that the interpreter-author can use
- hints are annotation in the interpreter
- they give information to the meta-tracer
- two hints presented here
- one to induce runtime feedback of arbitrary information
- the second one to influence constant folding

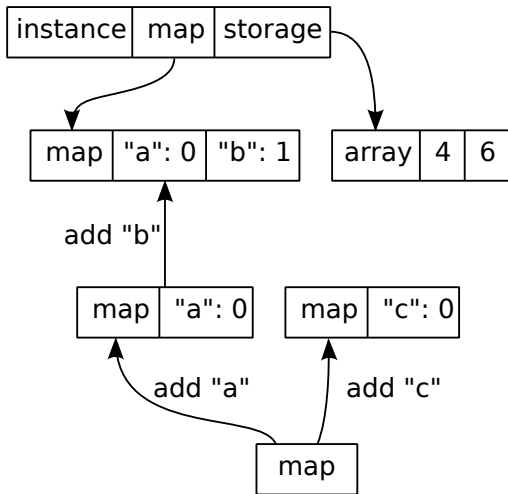
# Example: Instances with Maps

map	"a": 0	"b": 1
-----	--------	--------

# Example: Instances with Maps



# Example: Instances with Maps





# Map Implementation

```
class Map(object):
    def __init__(self, indexes):
        self.indexes = indexes
        ...

    def getindex(self, name):
        return self.indexes.get(name, -1)

    def add_attribute(self, name):
        ...

EMPTY_MAP = Map({})
```

```
class Instance(object):
    def __init__(self):
        self.map = EMPTY_MAP
        self.storage = []

    def getfield(self, name):
        index = self.map.getindex(name)
        if index != -1:
            return self.storage[index]
        return None

    def write_attribute(self, name, value):
        ...
```

## Trace for Code `inst.a + inst.b`

```
# inst1.getfield("a")  
map1 = inst1.map  
index1 = Map.getindex(map1, "a")  
guard(index1 != -1)  
storage1 = inst1.storage  
result1 = storage1[index1]
```

## Trace for Code `inst.a + inst.b`

```
# inst1.getfield("a")
map1 = inst1.map
index1 = Map.getindex(map1, "a")
guard(index1 != -1)
storage1 = inst1.storage
result1 = storage1[index1]

# inst1.getfield("b")
map2 = inst1.map
index2 = Map.getindex(map2, "b")
guard(index2 != -1)
storage2 = inst1.storage
result2 = storage2[index2]

v1 = result1 + result2
```

# Runtime Feedback Controlled by the Interpreter Author

- give the interpreter author a way to feed back runtime values into the trace
- written as `promote(x)`
- captures the argument's runtime value during tracing
- should be used only for variables that take few values

# Tiny Example

```
def f1(x, y):  
    promote(x)  
    z = x * 2 + 1  
    return z + y
```

```
guard(x1 == 4)  
v1 = x1 * 2  
z1 = v1 + 1  
v2 = z1 + y1  
return(v2)
```

# Foldable Operations Defined by the Interpreter Author

- let the interpreter author define foldable functions
- those functions typically don't look foldable
- otherwise there is no need for an annotation
- done via a function decorator `@elidable`

# Foldable Operations Defined by the Interpreter Author

- let the interpreter author define foldable functions
- those functions typically don't look foldable
- otherwise there is no need for an annotation
- done via a function decorator `@elidable`
- decorated functions should be pure
- or have idempotent side effects (such as a function that memoizes)
- trace optimizer will remove calls to such functions with constant arguments



# Adding Hints to Maps

```
class Map(object):  
    def __init__(self, indexes):  
        self.indexes = indexes  
        ...  
  
    @elidable  
    def getindex(self, name):  
        return self.indexes.get(name, -1)  
  
    def add_attribute(self, name):  
        ...  
  
EMPTY_MAP = Map({})
```

# Adding Hints to Maps

```
class Instance(object):
    def __init__(self):
        self.map = EMPTY_MAP
        self.storage = []

    def getfield(self, name):
        promote(self.map)
        index = self.map.getindex(name)
        if index != -1:
            return self.storage[index]
        return None

    def write_attribute(self, name, value):
        ...
```

## Trace with Hints for Code `inst.a + inst.b`

```
# inst1.getField("a")
map1 = inst1.map
guard(map1 == 0xb74af4a8)
index1 = Map.getIndex(map1, "a")
guard(index1 != -1)
storage1 = inst1.storage
result1 = storage1[index1]
```

```
# inst1.getField("b")
map2 = inst1.map
guard(map2 == 0xb74af4a8)
index2 = Map.getIndex(map2, "b")
guard(index2 != -1)
storage2 = inst1.storage
result2 = storage2[index2]
```

```
v1 = result1 + result2
```

```
map1 = inst1.map  
guard(map1 == 0xb74af4a8)  
storage1 = inst1.storage  
result1 = storage1[0]  
result2 = storage2[1]  
v1 = result1 + result2
```

# Uses of These Hints

`promote` lets one specialize on various things:

- user-level types
- shapes of instances
- the current state of a classes' methods
- ...

# Uses of These Hints

`promote` lets one specialize on various things:

- user-level types
- shapes of instances
- the current state of a classes' methods
- ...


uses of `@elidable`

- define immutable fields by decorating a getter
- declare arbitrary invariants

# Some Benchmarks

- benchmarks done using PyPy's Python interpreter
- about 30'000 lines of code
- 20 calls to `promote`
- 10 applications of `@elidable`

# Some Benchmarks





# Conclusion

- meta-tracing can make the efficient implementation of complex dynamic languages easier
- only requires to write a correct interpreter
- two kinds of hints to be added by the interpreter author allow arbitrary runtime feedback and its exploitation
- the hints are expressive enough to re-implement classical optimizations such as maps
- usage of the hints leads to good speedups for object-oriented code in PyPy's Python interpreter

# Bonus: Comparison with Partial Evaluation

# Bonus: Comparison with Partial Evaluation

- the only difference between meta-tracing and partial evaluation is that meta-tracing works

# Bonus: Comparison with Partial Evaluation

- the only difference between meta-tracing and partial evaluation is that meta-tracing works
- ... mostly kidding

# Bonus: Comparison with Partial Evaluation

- the only difference between meta-tracing and partial evaluation is that meta-tracing works
- ... mostly kidding
- very similar from the motivation and ideas
- PE was never scaled up to perform well on large interpreters
- classical PE mostly ahead of time
- PE tried very carefully to select the right paths to inline and optimize
- quite often this fails and inlines too much or too little
- tracing is much more pragmatic: simply look what happens