

A Way Forward in Parallelising Dynamic Languages

Position Paper

Armin Rigo

www.pypy.org

Remigius Meier


Department of Computer Science
ETH Zürich

Current situation

- Languages like Python & Ruby are very popular
- Concurrency using threads
 - e.g. background tasks
 - other models in other languages not covered here
- No parallelism in reference implementations
 - threads serialised using a single, global interpreter lock (**GIL**)
 - **no speedup** on multicore machines


The GIL

- Initially: single-threaded interpreters
- For concurrency, provide threads
 - not for parallelism
- **Challenges** to adapt interpreter:
 - reference counting GC, concurrent access to lists / dicts / objects
- **Easy solution:**
 - ultra coarse-grained locking
 - acquire GIL around the execution of bytecode instructions
 - atomic & isolated execution




Consequences —

- No parallelism
- GIL is interpreter-level sync, not available to applications
 - need application-level locking (semaphores, conditions, ...)
 - all challenges of concurrent programming still there



Consequences +

- Atomic & isolated instructions
 - things like `list.append()` are atomic
 - tons of websites mention this
 - latent races if language becomes really parallel —
- Sequential consistency
 - less surprises: “all variables volatile”
 - **global sequential order** of instructions



Sequential consistency

Example

Thread 1	Thread 2
A = B = 0	
A = 1 if B == 0:	B = 1 if A == 0:
At most one thread can enter here (seq. consistency)	

No sequential consistency:

→ both threads could enter (e.g. x86 CPUs)

Where do we want to go?

- Remove / replace the GIL
 - allow threads to run in parallel
- Keep GIL semantics
- Keep current APIs
 - no changes to existing concurrent / threaded apps required
- Then: find new ways
 - better synchronisation mechanism than locking
 - new models for new apps: AME, tuple spaces, actors, ...

Avoiding the GIL

Approaches

1. Fine-grained locking
2. Shared-nothing
3. Transactional memory

Comparison

- Performance
 - single threaded
 - parallelisation
- Backwards compatibility:
 - support for existing threaded applications
 - sequential consistency
 - atomic instructions (`list.append()`)
- Implementation effort (interpreter)
 - large open-source communities
- Bonus: better application-level synchronisation mechanism to replace locking
 - e.g. exposing interpreter-level synchronisation to application
 - better = composable, no deadlocks, scalable / parallelisable

1. Fine-grained locking

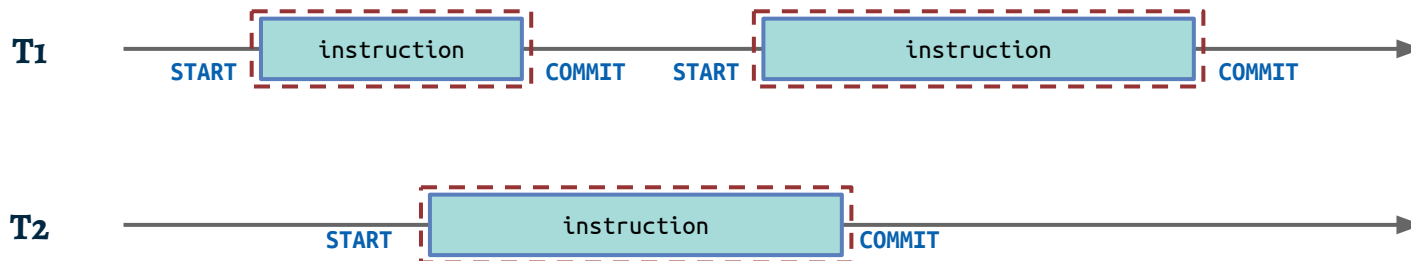
- Replace GIL with locks on objects / data structures
 - accessing different objects can run in parallel
- Possible to keep GIL semantics
- Harder to implement
 - many locks → deadlock risks
 - no centralised implementation
- Overhead of lock/unlock on objects
 - e.g. Jython depends on JVM for good lock removal
 - coarse GIL has less overhead
- Still needs application-level synchronisation

2. Shared-nothing

- Workaround instead of direct replacement
- Each independent part of the program gets its own interpreter (one GIL each)
- Simple implementation
- **Not compatible** with (most) existing threaded applications
 - extracting independent parts
- Explicit communication
 - good: clean model, no locks
 - bad: communication overhead

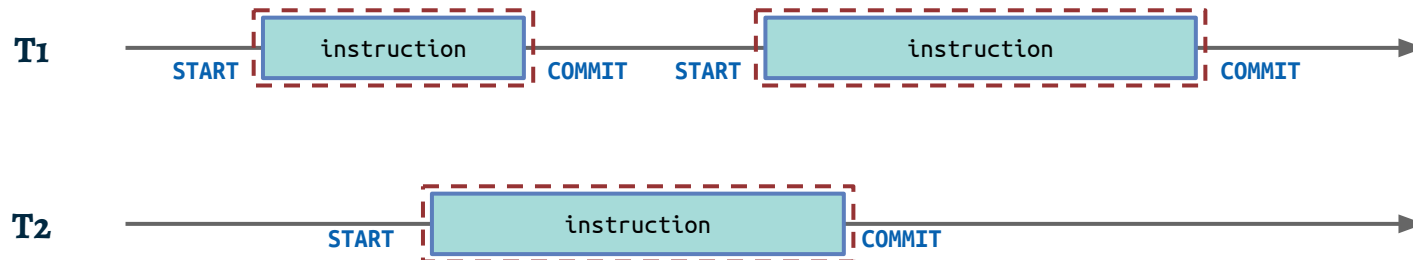
3. Transactional memory

- Transactions guarantee atomicity & isolation
- Direct replacement for the GIL
 - **lock** → transaction **START**
 - **unlock** → transaction **COMMIT**
- Keeps GIL semantics
 - sequential consistency (transaction schedule is serialisable)
 - instruction atomicity



3. Transactional memory

- Optimistically tries to execute in parallel
- Overhead:
 - validating memory accesses
 - start/commit/abort transactions
- Two categories:
 - **HTM**: hardware implementation
 - **STM**: software implementationalso: some hybrids



3.1 Hardware TM

- Easy, direct replacement with low overhead (<40%)
- Implementations not widespread
- **Limitations** (Intel Haswell):
 - false sharing on cache line level
 - transaction size ~ cache size
- Some transactions never succeed → needs fallback
 - in current approaches: GIL
 - our experiments suggest the fallback is needed often (limited scaling)
- Backwards compatible with existing applications
- Centralised implementation
- Still needs application-level synchronisation

3.2 Software TM

- Same benefits as HTM, except performance
 - often overhead of 100-1'000%
 - commonly scale well, but no speedup on <8 cores
- No hardware limits
- Unlimited transaction size
 - transactions can be exposed to the application as atomic blocks
 - not possible with HTM
- Atomic blocks
 - guarantee atomicity & isolation
 - composable, deadlock-free
 - move global locking protocol to the language implementation

Rough summary

	GIL	Fine-grained locking	Shared-nothing	HTM	STM
single-threaded performance	++	+	++	++	--
multi-threaded performance	--	+	+	++	++
backwards compatibility	++	++	--	++	++
implementation effort	++	-	++	++	++
synchronisation mechanism	o	o	+	o	++

- Picture not entirely clear
- Fine-grained locking & shared-nothing mixed
- HTM looks good
 - may never be widespread enough
 - hardware limitations in the current generation

Rough summary

- STM
 - synchronisation mechanism: atomic blocks
- Push for an easier parallel programming environment
 - sequential consistency & atomic blocks
 - going further than e.g. Java, C#, etc.
- Ultimately needs better performance
 - maybe the long-term solution?
 - hybrid TMs?
- Our direction of research:

Replace GIL with STM and move synchronisation from the application to the language implementation

Preview of PyPy-STM

- PyPy
 - Python interpreter
 - RPython: toolchain for producing dynamic language VMs
- Our own STM system: *STMGC-C7*
 - C library providing [STM](#) & [garbage collection](#)
 - optimised for use in dynamic language VMs (to replace GIL)
 - tight TM-GC cooperation: shared barriers, object lifetime optimisations
 - [low overhead](#)

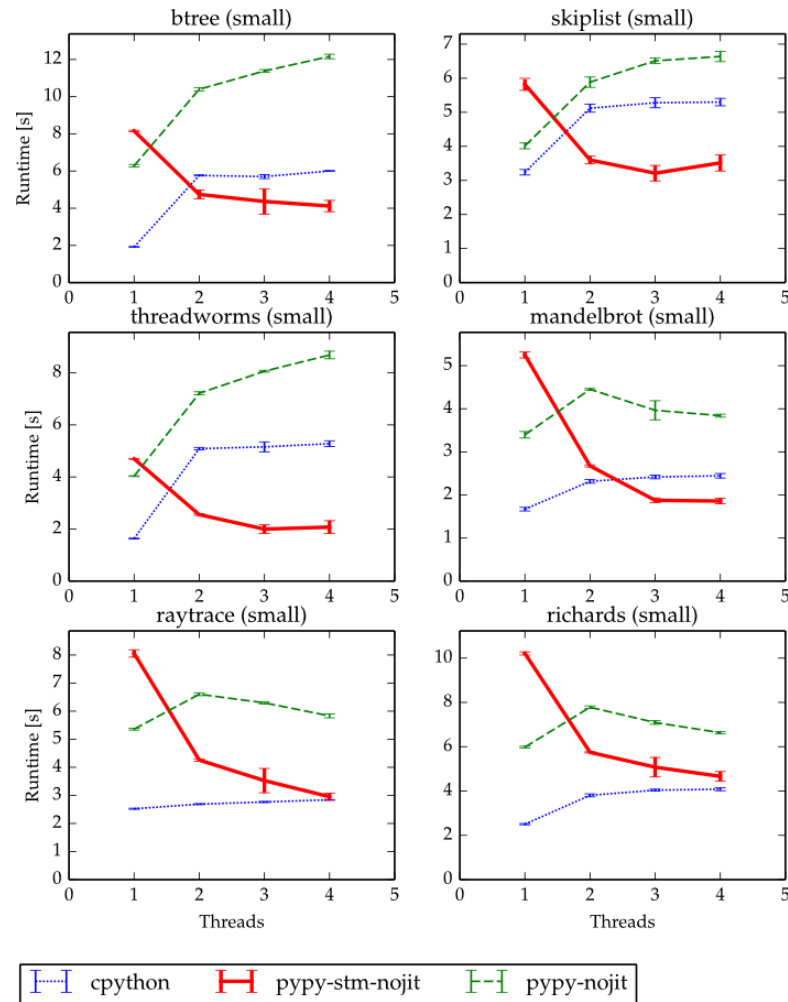
Preview of PyPy-STM

Single Threaded

- Max. 100% overhead
- Avg. 45% overhead

Multithreaded

- Speedup 1.1 - 1.9×



Questions & Discussion