

Software Transactional Memory on PyPy

Pseudo-Goal

- ▶ “Kill the GIL”
- ▶ GIL = Global Interpreter Lock

Real Goals

- ▶ Multi-core programming
- ▶ But *reasonable* multi-core programming
- ▶ Using the recent model of Transactional Memory

PyPy-STM

- ▶ An executable `pypy-stm` which uses internally Software Transactional Memory
- ▶ Optimistically run multiple threads in parallel
- ▶ The only new feature is `atomic`:

```
with atomic:  
    piece of code...
```

Example of higher-level API

```
def work (...):  
    ...  
    several more calls to:  
        transaction.add(work, ...)  
    ...
```

- ▶ Starts N threads, scheduling *work()* calls to them
- ▶ Each *work()* is done in an `atomic` block
- ▶ Multi-core, but as if all the *work()* are done sequentially

Status

- ▶ Kind of working without the JIT
- ▶ Roughly three times slower (you need four cores to see benefits)
- ▶ Working on the JIT support

Q&A

- ▶ Thank you!
- ▶ Budget of \$10k left, likely more needed too