

Allocation Removal by Partial Evaluation in a Tracing JIT

Carl Friedrich Bolz¹ Antonio Cuni¹ Maciej Fijałkowski²
Michael Leuschel¹ Samuele Pedroni³ Armin Rigo¹

¹Heinrich-Heine-Universität Düsseldorf, STUPS Group, Germany

²merlinux GmbH, Hildesheim, Germany

³Open End, Göteborg, Sweden

2011 Workshop on Partial Evaluation and Program
Manipulation, January 24, 2011

Dynamic Languages are Slow

- Interpretation overhead
- Type dispatching
- Boxing of primitive types

Dynamic Languages are Slow

- Interpretation overhead
- Type dispatching
- Boxing of primitive types
 - A lot of allocation of short-lived objects

Dynamic Languages are Slow: Example

$x = a + b$; $y = x + c$
evaluated in an interpreter:

Dynamic Languages are Slow: Example

$x = a + b$; $y = x + c$

evaluated in an interpreter:

- 1 What's the type of a ? Integer
- 2 What's the type of b ? Integer

Dynamic Languages are Slow: Example

$x = a + b; y = x + c$

evaluated in an interpreter:

- ① What's the type of a? Integer
- ② What's the type of b? Integer
- ③ unbox a
- ④ unbox b
- ⑤ compute the sum

Dynamic Languages are Slow: Example

$x = a + b; y = x + c$

evaluated in an interpreter:

- ① What's the type of a? Integer
- ② What's the type of b? Integer
- ③ unbox a
- ④ unbox b
- ⑤ compute the sum
- ⑥ box the result as an Integer
- ⑦ store into x

Dynamic Languages are Slow: Example

$x = a + b; y = x + c$

evaluated in an interpreter:

- 1 What's the type of a? Integer
- 2 What's the type of b? Integer
- 3 unbox a
- 4 unbox b
- 5 compute the sum
- 6 box the result as an Integer
- 7 store into x
- 8 What's the type of x? Integer
- 9 What's the type of c? Integer

Dynamic Languages are Slow: Example

$x = a + b; y = x + c$

evaluated in an interpreter:

- 1 What's the type of a? Integer
- 2 What's the type of b? Integer
- 3 unbox a
- 4 unbox b
- 5 compute the sum
- 6 box the result as an Integer
- 7 store into x
- 8 What's the type of x? Integer
- 9 What's the type of c? Integer
- 10 unbox x
- 11 unbox c
- 12 compute the sum

Dynamic Languages are Slow: Example

$x = a + b; y = x + c$

evaluated in an interpreter:

- 1 What's the type of a? Integer
- 2 What's the type of b? Integer
- 3 unbox a
- 4 unbox b
- 5 compute the sum
- 6 box the result as an Integer
- 7 store into x
- 8 What's the type of x? Integer
- 9 What's the type of c? Integer
- 10 unbox x
- 11 unbox c
- 12 compute the sum
- 13 box the result as an Integer
- 14 store into y

What to do?

- Use a JIT compiler
- **Add an optimization that can deal with heap operations**

What to do?

- Use a JIT compiler
- **Add an optimization that can deal with heap operations**
- optimize short-lived objects
- remove some of the redundant type checks

Overview

- 1 Experimental Context
- 2 Proposed Optimization
- 3 Benchmarks

Our Experimental Context is the PyPy Project

A general environment for implementing dynamic languages

Our Experimental Context is the PyPy Project

A general environment for implementing dynamic languages

Approach

- write an interpreter for the language in RPython
- compilable to an efficient C-based VM

Our Experimental Context is the PyPy Project

A general environment for implementing dynamic languages

Approach

- write an interpreter for the language in RPython
- compilable to an efficient C-based VM
- (RPython is a restricted subset of Python)

PyPy Uses a Novel Tracing JIT

the feature that makes PyPy interesting:

- a meta-JIT, applicable to many languages
- needs a few source-code hints (or user annotations) in the interpreter
- JIT is a tracing JIT compiler

Tracing JITs Compile by Observing an Interpreter

- VM contains both an interpreter and the tracing JIT compiler
- JIT works by observing and logging what the interpreter does
 - for interesting, commonly executed code paths
 - produces a linear list of operations (trace)
- trace is optimized and then turned into machine code

The Advantages of Tracing JITs

- Traces are interesting linear pieces of code
- most of the time correspond to loops
- everything called in the trace is inlined
- can perform good optimizations on the trace
- rarer paths run by the interpreter

Example Trace

Trace of $x = a + b$; $y = x + c$:

```
guard_class(a, Integer)
guard_class(b, Integer)
i1 = get(a, intval)
i2 = get(b, intval)
i3 = int_add(i1, i2)
x = new(Integer)
set(x, intval, i3)
```

Example Trace

Trace of $x = a + b$; $y = x + c$:

```
guard_class(a, Integer)
```

```
guard_class(b, Integer)
```

```
i1 = get(a, intval)
```

```
i2 = get(b, intval)
```

```
i3 = int_add(i1, i2)
```

```
x = new(Integer)
```

```
set(x, intval, i3)
```

```
guard_class(x, Integer)
```

```
guard_class(c, Integer)
```

```
i4 = get(x, intval)
```

```
i5 = get(c, intval)
```

```
i6 = int_add(i4, i5)
```

```
y = new(Integer)
```

```
set(y, intval, i6)
```

```
return(y)
```

Optimized Example Trace

Trace of $x = a + b$; $y = x + c$:

```
guard_class(a, Integer)
```

```
guard_class(b, Integer)
```

```
i1 = get(a, intval)
```

```
i2 = get(b, intval)
```

```
i3 = int_add(i1, i2)
```

```
x = new(Integer)
```

```
set(x, intval, i3)
```

```
guard_class(x, Integer)
```

```
guard_class(c, Integer)
```

```
i4 = get(x, intval)
```

```
i5 = get(c, intval)
```

```
i6 = int_add(i3, i5)
```

```
y = new(Integer)
```

```
set(y, intval, i6)
```

```
return(y)
```

Contribution of our Paper

- a simple, efficient and effective optimization of heap operations in a trace
- using online partial evaluation
- fully implemented and in use in large-scale interpreters

Contribution of our Paper

- a simple, efficient and effective optimization of heap operations in a trace
- using online partial evaluation
- fully implemented and in use in large-scale interpreters

Ingredients


- a slightly simplified model for objects on the heap
- operational semantics of trace operations that manipulate the heap
- optimization rules for those operations, following the operational semantics

Heap Model


0x01

T1		
----	--	--

Heap Model



Heap Model




Operations Manipulating Heap Objects

- `v = new(T)` makes a new object
- `u = get(w, F)` reads a field out of an object
- `set(v, F, w)` writes a field of an object
- `guard(v, T)` checks the type of an object


Operations: New

`v=new(T4)`




Operations: New

v=new(T4)




Operations: Get

$v = \text{get}(w, L)$




Operations: Get

$v = \text{get}(w, L)$




Operations: Set

set(w, R, u)




Operations: Set

set(w, R, u)




Operations: Guard

guard(u, T1)




Operations: Guard

guard(u, T1)




Operations: Guard

guard(w, T1)



Operations: Guard

guard(w, T1)



Optimization by Online Partial Evaluation


- Trace is optimized using online partial evaluation
- part of the runtime heap is modelled in the static heap
- static heap contains objects that are allocated within the trace
- (as opposed to before the trace is executed)

Optimization by Online Partial Evaluation

- Trace is optimized using online partial evaluation
- part of the runtime heap is modelled in the static heap
- static heap contains objects that are allocated within the trace
- (as opposed to before the trace is executed)
- operations acting on the static heap can be executed
- follows operational semantics
- all others need to be residualized
- all fairly straightforward

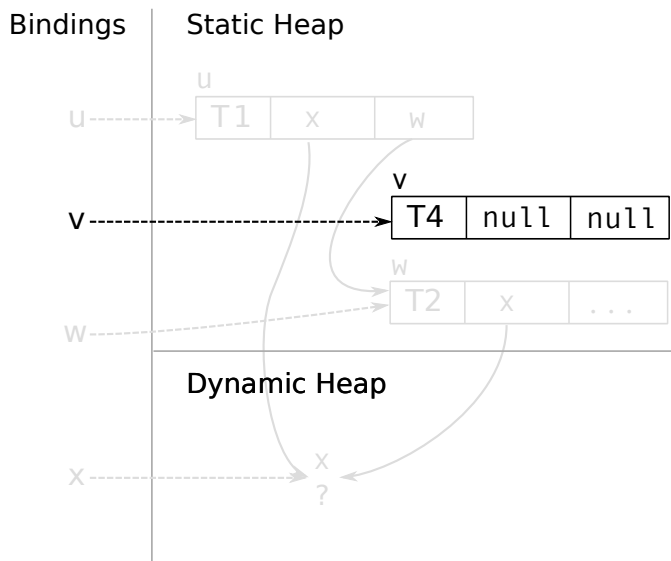
Optimizing New

`v=new(T4)`




Optimizing New

~~v=new(T4)~~




Optimizing Get

$v = \text{get}(w, L)$




Optimizing Get

~~$v = \text{get}(w, L)$~~




Optimizing Get

$v = \text{get}(x, R)$




Optimizing Get




Optimizing Guard

guard(u, T1)




Optimizing Guard

~~guard(u, T1)~~




Optimizing Guard

guard(v, T1)




Optimizing Guard




Optimizing Set

set(u,R,v)



Optimizing Set

~~set(u,R,v)~~



Problem: What happens if we write a static object into a dynamic one?

Problem: What happens if we write a static object into a dynamic one?


- lose track of the static object because of possibility of aliasing
- need to lift the static object
- lifting produces operations that recreate the static object

Problem: What happens if we write a static object into a dynamic one?

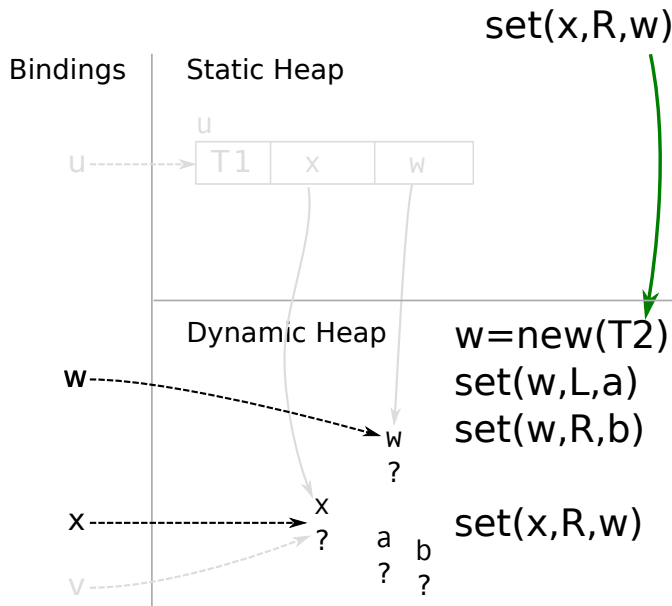
- lose track of the static object because of possibility of aliasing
- need to lift the static object
- lifting produces operations that recreate the static object
- needs to be careful due to recursive structures

Optimizing Set

set(x,R,w)



Optimizing Set



Properties of the Optimization

- output trace is never longer than input trace
- runtime linear in the length of the trace

Properties of the Optimization

- output trace is never longer than input trace
- runtime linear in the length of the trace

Implementation


- about 400 lines of code
- some added complexity over presentation
- objects with arbitrary numbers of fields
- array support

Optimizing the Example Trace

Bindings

Static Heap

Dynamic Heap




```
guard(a, Int)
guard(b, Int)
i1 = get(a, intval)
i2 = get(b, intval)
i3 = int_add(i1, i2)
x = new(Int)
set(x, intval, i3)
guard(x, Int)
guard(c, Int)
i4 = get(x, intval)
i5 = get(c, intval)
i6 = int_add(i4, i5)
y = new(Int)
set(y, intval, i6)
return(y)
```

Optimizing the Example Trace

Bindings

Static Heap

Dynamic Heap



```
guard(a, Int)
```

```
guard(b, Int)
```

```
i1 = get(a, intval)
```

```
i2 = get(b, intval)
```

```
i3 = int_add(i1, i2)
```

```
x = new(Int)
```

```
set(x, intval, i3)
```

```
guard(x, Int)
```

```
guard(c, Int)
```

```
i4 = get(x, intval)
```

```
i5 = get(c, intval)
```

```
i6 = int_add(i4, i5)
```

```
y = new(Int)
```

```
set(y, intval, i6)
```


```
return(y)
```

Optimizing the Example Trace

Bindings

Static Heap

Dynamic Heap



```
guard(a, Int)
```

```
guard(b, Int)
```

```
i1 = get(a, intval)
```

```
i2 = get(b, intval)
```

```
i3 = int_add(i1, i2)
```

```
x = new(Int)
```

```
set(x, intval, i3)
```

```
guard(x, Int)
```

```
guard(c, Int)
```

```
i4 = get(x, intval)
```

```
i5 = get(c, intval)
```

```
i6 = int_add(i4, i5)
```

```
y = new(Int)
```

```
set(y, intval, i6)
```


```
return(y)
```

Optimizing the Example Trace

Bindings

Static Heap

Dynamic Heap




```
guard(a, Int)
guard(b, Int)
i1 = get(a, intval)
i2 = get(b, intval)
i3 = int_add(i1, i2)
x = new(Int)
set(x, intval, i3)
guard(x, Int)
guard(c, Int)
i4 = get(x, intval)
i5 = get(c, intval)
i6 = int_add(i4, i5)
y = new(Int)
set(y, intval, i6)
return(y)
```

Optimizing the Example Trace

Bindings

Static Heap

Dynamic Heap




```
guard(a, Int)
guard(b, Int)
i1 = get(a, intval)
i2 = get(b, intval)
i3 = int_add(i1, i2)
x = new(Int)
set(x, intval, i3)
guard(x, Int)
guard(c, Int)
i4 = get(x, intval)
i5 = get(c, intval)
i6 = int_add(i4, i5)
y = new(Int)
set(y, intval, i6)
return(y)
```


Optimizing the Example Trace

Bindings

Static Heap

Dynamic Heap




```
guard(a, Int)
guard(b, Int)
i1 = get(a, intval)
i2 = get(b, intval)
i3 = int_add(i1, i2)
x = new(Int)
set(x, intval, i3)
guard(x, Int)
guard(c, Int)
i4 = get(x, intval)
i5 = get(c, intval)
i6 = int_add(i4, i5)
y = new(Int)
set(y, intval, i6)
return(y)
```

Optimizing the Example Trace

Bindings

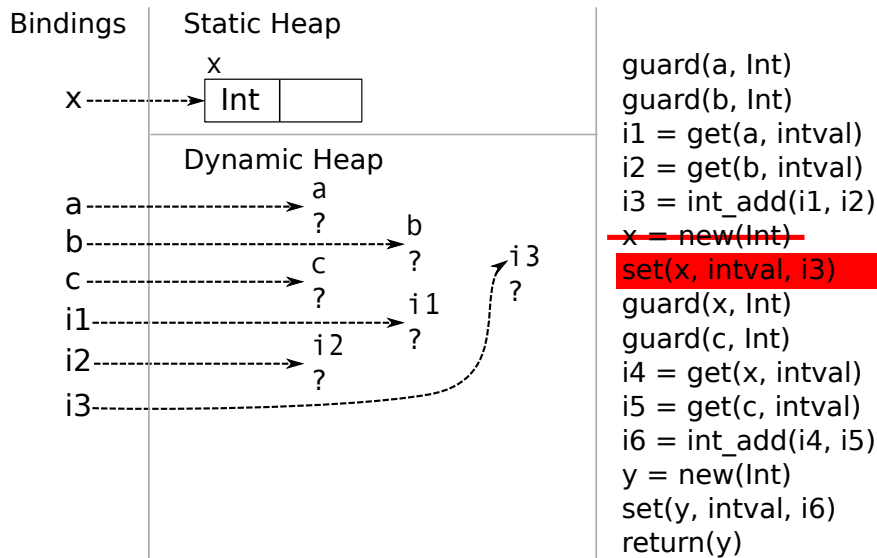
Static Heap

Dynamic Heap

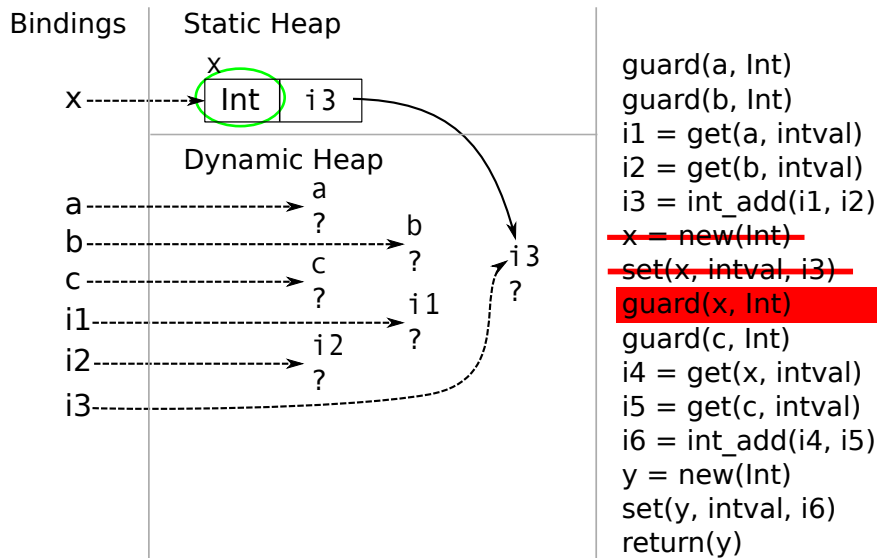


```
guard(a, Int)
guard(b, Int)
i1 = get(a, intval)
i2 = get(b, intval)
i3 = int_add(i1, i2)
x = new(Int)
set(x, intval, i3)
guard(x, Int)
guard(c, Int)
i4 = get(x, intval)
i5 = get(c, intval)
i6 = int_add(i4, i5)
y = new(Int)
set(y, intval, i6)
return(y)
```

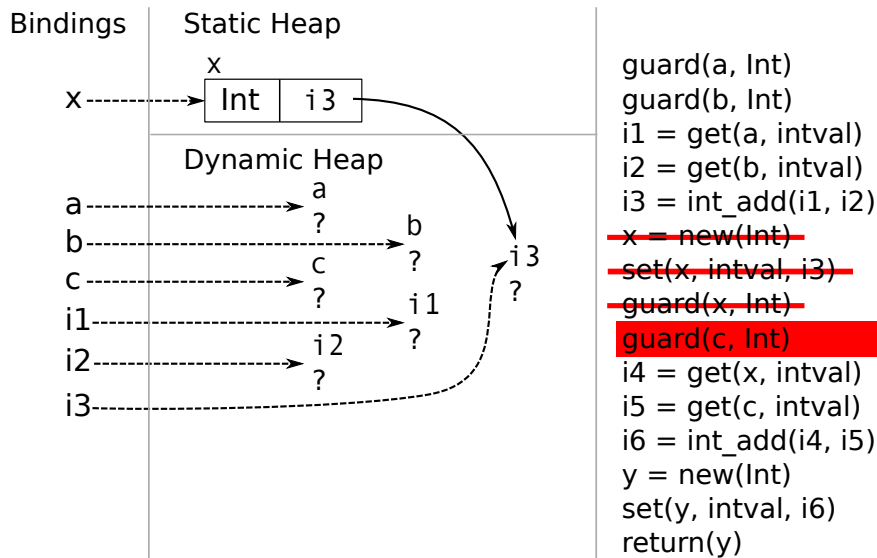
Optimizing the Example Trace



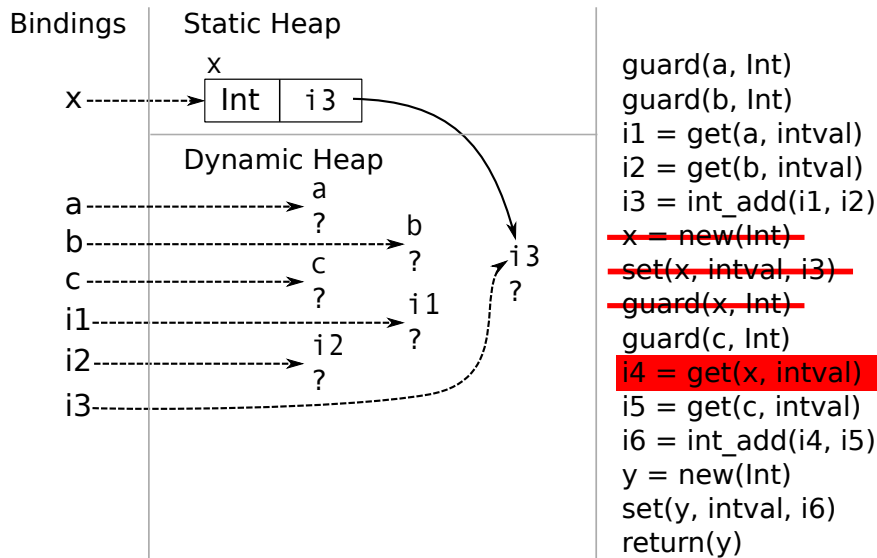
Optimizing the Example Trace



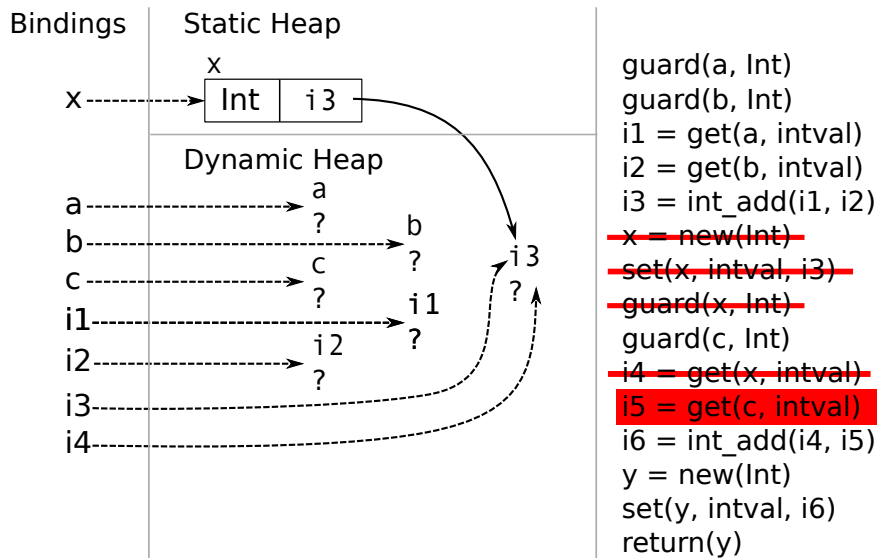
Optimizing the Example Trace



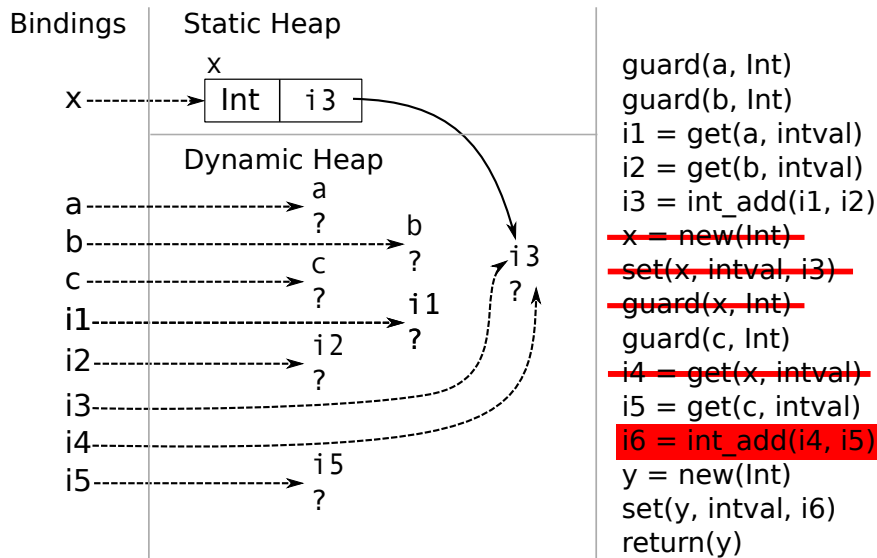
Optimizing the Example Trace



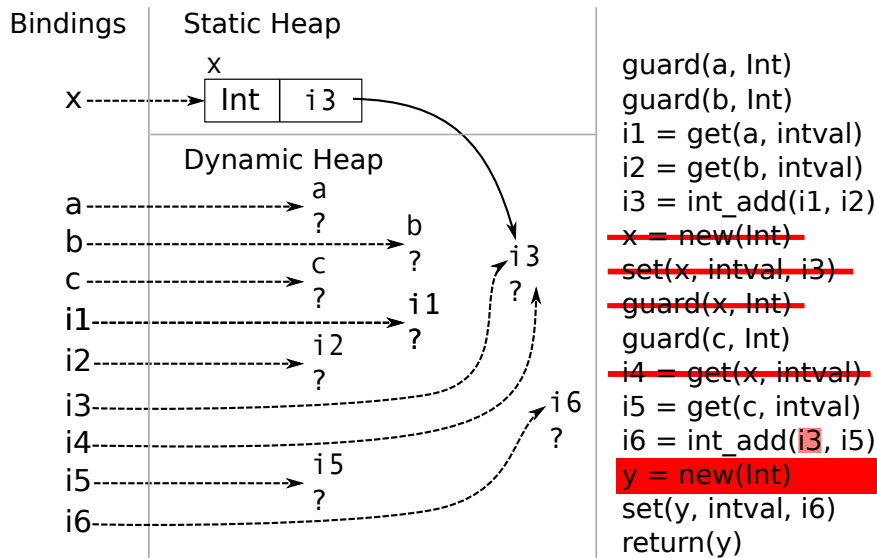
Optimizing the Example Trace



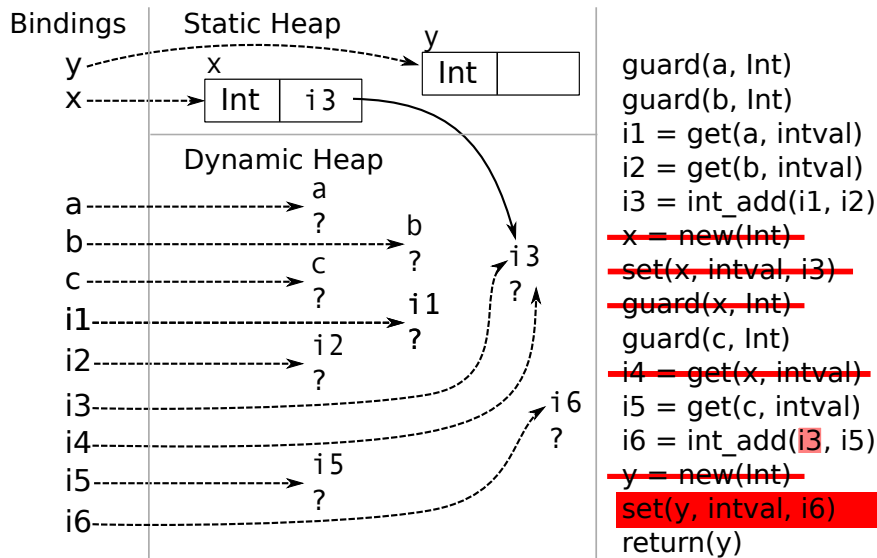
Optimizing the Example Trace




Optimizing the Example Trace



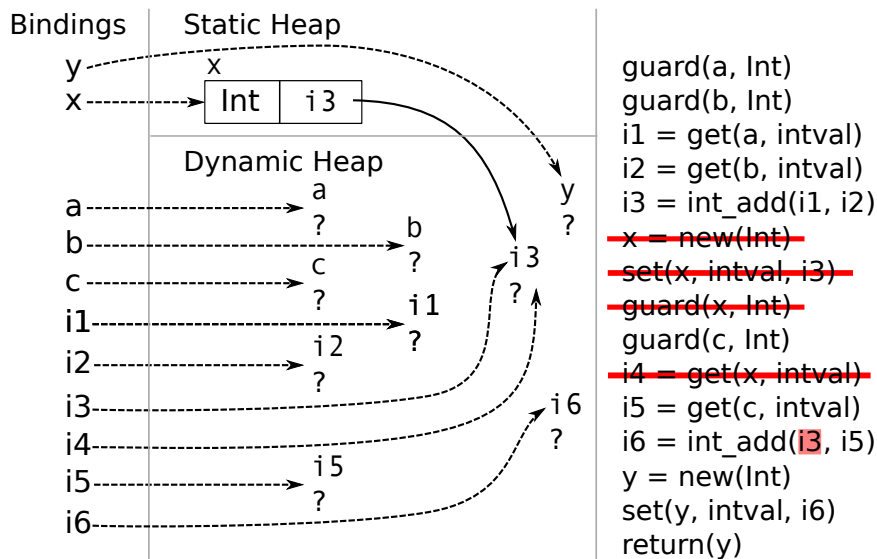
Optimizing the Example Trace



Optimizing the Example Trace



Optimizing the Example Trace



Benchmark Results

- to evaluate the optimization we used PyPy's Python interpreter with real-world programs
 - interpreter is about 30'000 lines of code


Benchmark Results

- to evaluate the optimization we used PyPy's Python interpreter with real-world programs
 - interpreter is about 30'000 lines of code
- optimization can remove
 - 70% of all new operations
 - 14% of all get/set operations
 - 93% of all guard operations

Benchmark Results

- to evaluate the optimization we used PyPy's Python interpreter with real-world programs
 - interpreter is about 30'000 lines of code
- optimization can remove
 - 70% of all new operations
 - 14% of all get/set operations
 - 93% of all guard operations
- Timings improve by a factor between 1.1 and 6.95
- outperforming standard Python on all benchmarks but two

Benchmark



Conclusion

- We propose a very simple partial-evaluation-based optimization for tracing JITs of dynamic languages that:
 - can remove a lot of allocations and type checks in practical programs.
 - is efficient and effective.
 - has no control issues because all control decisions are made by the tracing JIT.

Conclusion

- We propose a very simple partial-evaluation-based optimization for tracing JITs of dynamic languages that:
 - can remove a lot of allocations and type checks in practical programs.
 - is efficient and effective.
 - has no control issues because all control decisions are made by the tracing JIT.
- We claim that this is a general strategy to get rid of control problems by simply observing the runtime behaviour of the program.

Backup Slides

What About Correctness?

- We haven't proven correctness yet
- should not be too hard
- lifting needs to be carefully handled

Comparison to Escape Analysis

- Effect very similar to escape analysis
- Escape analysis needs a complex upfront analysis
- our optimization automatically has a lot of context, due to the inlining tracing does
- our optimization can optimize operations on objects even if they escape later

Comparison to "Dynamic Typing" /Boxing Analysis

- those optimizations work ahead of time
- don't work for many dynamic languages, where the source simply does not contain enough information

Python Example:

```
def sum(container, initial):  
    result = initial  
    for element in container:  
        result = result + element  
    return result
```