



# **OPTIMIZING PYTHON PROGRAMS, PYPY TO RESCUE**

6. Oct. 2016, Cape Town

**RICHARD PLANGGER**

# MORE "GENERAL" PYPY TALK

Goals:

- An approach to optimize Python programs
- Examples
- How not to start optimizing
- What is PyPy up to now?

# PYPY IS A ...

**... fast virtual machine for Python**

developed by researchers, freelancers and many contributors.

```
$ python yourprogram.py
```

```
$ pypy yourprogram.py
```

**PYPY IS NOT JUST  
THAT**

Experiment with new ideas

- Python written in Python
- RPython
- JIT compiler
- VMProf
- PyPy STM
- ...

# ABOUT ME

Working on PyPy (+1,5y)

Master thesis → GSoC 2015 → PyPy

living and working in Austria



# **SPEEDY PYTHON PROGRAMS?**

When is your Python program fast enough?

When it gets a speeding ticket because it is too fast?

or when PyPy's benchmark suite reaches 10x faster on average?

**Neither**

Run your program and measure your **criteria**

# FOR EXAMPLE?

- CPU time
- Peak Heap Memory
- Requests per second
- Latency
- ...

Dissatisfaction with one criteria of your program!

**SOME THEORY ...**

# COMPLEXITY

Big-O-Notation

Classify e.g. a function and it's processing time

Increase input size to the function



- `a = 3` #  $O(1)$
- `[x+1 for x in range(n)]` #  $O(n)$
- `[[x+y for x in range(n)] \`  
`for y in range(m)]` #  $O(n*m) == O(n)$  if  $n > m$

Bubble sort vs Quick Sort

$O(n^2)$  vs  $O(n \log n)$

# COMPLEXITY

Yields the most gain, independent from the language

E.g. prefer  $O(n)$  over  $O(n^{**2})$

# ONLY OPTIMIZE A ROUTINE IF ...

you know that the complexity cannot be stripped down

# LET'S START FROM THE BEGINNING

with a small example

# READING LOG FILES!

JITLOG (facility to observe PyPy's JIT internals)

- Written in Python
- Moved to [vmprof.com](http://vmprof.com)
- Log files can easily take up to 40MB uncompressed
- Takes ~10 seconds to parse with CPython
- Complexity is linear to input size of the log file

# THANKS TO PYTHON

+ Little development time

+ Easy to test



- Takes too long to parse
- Parsing is done each request

Our criteria: CPU time too long + requests per second  
(Many objects are allocated)

# SUGGESTION

Caching

Reduce CPU time

Let's have both

Caching - Easily done with your favourite caching framework

Reduce CPU time - PyPy seems to be good at that?

# LET'S RUN IT...

```
$ cpython2.7 parse.py 40mb.log  
~ 10 seconds
```

```
$ pypy2 parse.py 40mb.log  
~ 2 seconds
```

# CACHING

Requests really feel instant after the log has been loaded  
once

Precache

# **THE LAZY APPROACH OF OPTIMIZING PYTHON**

# VMPROF

```
$ pip install vmprof
```

```
$ python -m vmprof --web parse.py
```

→ [link](#)

(100.00%)

00% ↓ 0.00% / 0.00%

JIT: 57.74%

GC: 0.00%



# INTRODUCING PYPY'S

<module>

main

JIT

\_parse\_jitlog

\_parse\_jitlog

read\_resop\_descr

read\_merge\_point

add\_instr

<dictcomp>



# HOT SPOTS

Loops / Repeat construct!

What kind program can you build without loops?

# A SIMPLIFIED VIEW


1. Start interpretation
2. Loops trigger recording
3. Optimization stage
4. Machine code generation

# BEYOND THE SCOPE OF LOOPS


Guards ensure correctness

Frequent guard failure triggers recording

# PERCEPTION



<http://abstrusegoose.com/secretarchives/under-the-hood> - CC BY-NC 3.0 US



Traced Opt **Rewritten** Asm

→ link  
This trace is a **loop**. Driver name: "" Scope/Function: "schedule"  
Filename: ../pypy/rpython/translator/goal/richards.py Lineno: 351

```
1: increment_debug_counter(9053872) passed 980.8 k times
2: gc_load_i(p0, 40, -8)
3: gc_load_r(p0, 64, 8)
4: gc_load_i(p0, 72, -8)
5: gc_load_r(p0, 16, 8)
6: gc_load_r(p0, 32, 8)
7: gc_load_r(p0, 48, 8)
8: gc_load_r(p0, 56, 8)
9: gc_load_r(p24, 16, 8)
10: gc_load_r(p24, 24, 8)
11: gc_load_r(p24, 32, 8)
12: gc_load_r(p24, 40, 8)
```

# JITVIEWER

Tool to inspect PyPy internals

Helps you to learn and understand PyPy

Provided at [vmprof.com](http://vmprof.com)

# PROPERTIES & TRICKS

- Type specialization
- Object unboxing
- GC scheme
- Dicts
- Dynamic class creation (Instance maps)
- Function calls (+ Inlining)

# **ANOTHER REAL WORLD EXAMPLE**



# MAGNETIC

Marketing tech company

Switched to PyPy 3 years ago

# **Q: WHAT DOES YOUR SERVICE DO?**

A: ... allow generally large companies to send targeted marketing (e.g. serve ads) to people based on data we have learned

**Q: PYPY, WHERE WAS IT  
MOST HELPFUL?**

A: ... ~30% speedups immediately from switching to PyPy ...

# Q: PYPY ISSUES?

A: ... we had to solve for rolling deploys ... but that's ok,  
that's fairly easy ...

# Q: VALUE TO YOUR COMPANY?

A: Latency speedup was somewhere around 10% ...

But that number is deceiving

It's very valuable for us obviously

But it's only 10%, because even this app that I'm talking about, which is fairly high volume (500,000 QPS), is a WSGI app

So it spends lots of time blocking

# TIMEIT

why not use perf?

Try timeit on PyPy

# PYTHON 3.5

Progressed quite a bit

`async io`

Many more small details (sprint?)

# C-EXTENSIONS

NumPy on top of the emulated layer

Boils down to managing PyPy & CPython objects



# CLOSING EXAMPLE

how to move from cpu limited to network limited

[link](#)

**QUESTIONS?**

[morepypy.blogspot.com](http://morepypy.blogspot.com)

[software@vimloc.systems](mailto:software@vimloc.systems)

Join on IRC [#pypy](#)