

Multi-core programming without threads: the transactional memory approach

Armin Rigo and Remi Meier

A common story

- Write a program or website
- The program gets popular
- Datasets grow bigger and bigger

→ The program needs to scale

Problem

“We have all these cores, we really need to use more than one now”

- Multiple processes? Not if the program is too “irregular”
- Multiple threads?

Threads are hard


- Concurrent, independent execution streams
- Shared memory
 - Manual synchronization and coordination
 - Locks: Locks are hard

Our approach

- Not automatic parallelization, but automatic synchronization
- Programmer...
 - identifies code where parallel execution is ***possible***
 - splits code into atomic units that may run in any order
- Runtime...
 - ensures atomic and isolated execution
 - ensures serializable schedule
 - handles conflicts transparently

Example

```
...  
for block in pending:  
    blockcount += 1  
    transform(block)  
    already_seen.add(block)
```



```
...  
tq = transaction.TransactionQueue()  
for block in pending:  
    tq.add(transform, block)  
tq.run()  
blockcount += len(pending)  
already_seen.update(pending)
```


- The calls to `transform()` should be mostly independent
- But they do not need to be *fully* independent
- Benefits from parallelism **if possible**

Our solution

- Transactional memory (TM)
 - coarse-grained atomicity and serializability
- Modified Python VM (PyPy)
- All Python-level code executes in TM
 - no statically known code boundaries for “parallel units”
- Completely transparent / native:
 - TM in the virtual machine
 - no explicit, exposed TM
 - only “TransactionQueue”, an API for running parallel units

Speculative parallel execution

- Application enqueues units
- Units get scheduled
- Unit runs fully in one transaction
 - atomically, isolated
 - in parallel if possible
- Units schedule is serializable



Consequences

- No (exposed) threads, no manual synchronization
- Gradually introduce parallelism
 - programmer marks promising spots
 - if parallel execution is not possible, the code still executes correctly...
 - ...but not faster: programmer may have to tweak the code to minimize conflicts
- Requires execution of **all** code in TM
 - performance challenge
 - so far, 25%-40% typical performance hit, up to 100% in some cases (but fully working :-)
 - Just-In-Time compiler may help more


STMGC

- C library that handles both GC and STM
- Used by PyPy
 - benefits from PyPy being written in “RPython”, a flexible language we control
 - but can be used more generally

STMGC API

- Classical part: the GC
 - `obj = allocate(size)`
 - pushing/popping on the “root stack” of objects
 - finding pointers inside objects
- New part: STM
 - `stm_read(obj)` `/* no return value, extremely cheap */`
 - `stm_write(obj)` `/* no return value */`
 - `stm_start_transaction()`, `stm_commit_transaction()`
- Concurrent transactions see separate memory
- Unmodified memory pages are shared (with `mmap`)

Memory views



Conclusion

- Not automatic parallelization, but close
- Works on existing, irregular programs
 - No need to learn a specific model
- Drawback: the programmer may have to add some tweaks to reduce conflicts
 - But the program is always “correct”
- STMGC+PyPy shows a reasonable(?) performance hit so far